

```

import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import warnings
warnings.filterwarnings('ignore')

df=pd.read_csv("heart.csv")

df.head(5)

   Age Sex ChestPainType RestingBP Cholesterol FastingBS RestingECG
MaxHR \
0   40   M             ATA       140        289         0      Normal
172
1   49   F             NAP       160        180         0      Normal
156
2   37   M             ATA       130        283         0          ST
98
3   48   F             ASY       138        214         0      Normal
108
4   54   M             NAP       150        195         0      Normal
122

   ExerciseAngina Oldpeak ST_Slope HeartDisease
0              N     0.0      Up           0
1              N     1.0      Flat          1
2              N     0.0      Up           0
3              Y     1.5      Flat          1
4              N     0.0      Up           0

df.shape
(918, 12)

df.isna().sum()

Age          0
Sex          0
ChestPainType 0
RestingBP     0
Cholesterol   0
FastingBS     0
RestingECG    0
MaxHR         0
ExerciseAngina 0
Oldpeak        0
ST_Slope       0
HeartDisease   0
dtype: int64

```

```

df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 918 entries, 0 to 917
Data columns (total 12 columns):
 #   Column            Non-Null Count  Dtype  
--- 
 0   Age               918 non-null    int64  
 1   Sex               918 non-null    object  
 2   ChestPainType    918 non-null    object  
 3   RestingBP         918 non-null    int64  
 4   Cholesterol      918 non-null    int64  
 5   FastingBS        918 non-null    int64  
 6   RestingECG       918 non-null    object  
 7   MaxHR             918 non-null    int64  
 8   ExerciseAngina   918 non-null    object  
 9   Oldpeak           918 non-null    float64 
 10  ST_Slope          918 non-null    object  
 11  HeartDisease     918 non-null    int64  
dtypes: float64(1), int64(6), object(5)
memory usage: 86.2+ KB

df.describe()

          Age      RestingBP      Cholesterol      FastingBS      MaxHR \ 
count  918.000000  918.000000  918.000000  918.000000  918.000000
mean   53.510893  132.396514  198.799564  0.233115   136.809368
std    9.432617   18.514154   109.384145  0.423046   25.460334
min    28.000000   0.000000   0.000000   0.000000   60.000000
25%   47.000000   120.000000  173.250000  0.000000   120.000000
50%   54.000000   130.000000  223.000000  0.000000   138.000000
75%   60.000000   140.000000  267.000000  0.000000   156.000000
max   77.000000   200.000000  603.000000  1.000000   202.000000

          Oldpeak      HeartDisease \ 
count  918.000000  918.000000
mean   0.887364   0.553377
std    1.066570   0.497414
min   -2.600000   0.000000
25%   0.000000   0.000000
50%   0.600000   1.000000
75%   1.500000   1.000000
max   6.200000   1.000000

df['Sex'].unique()

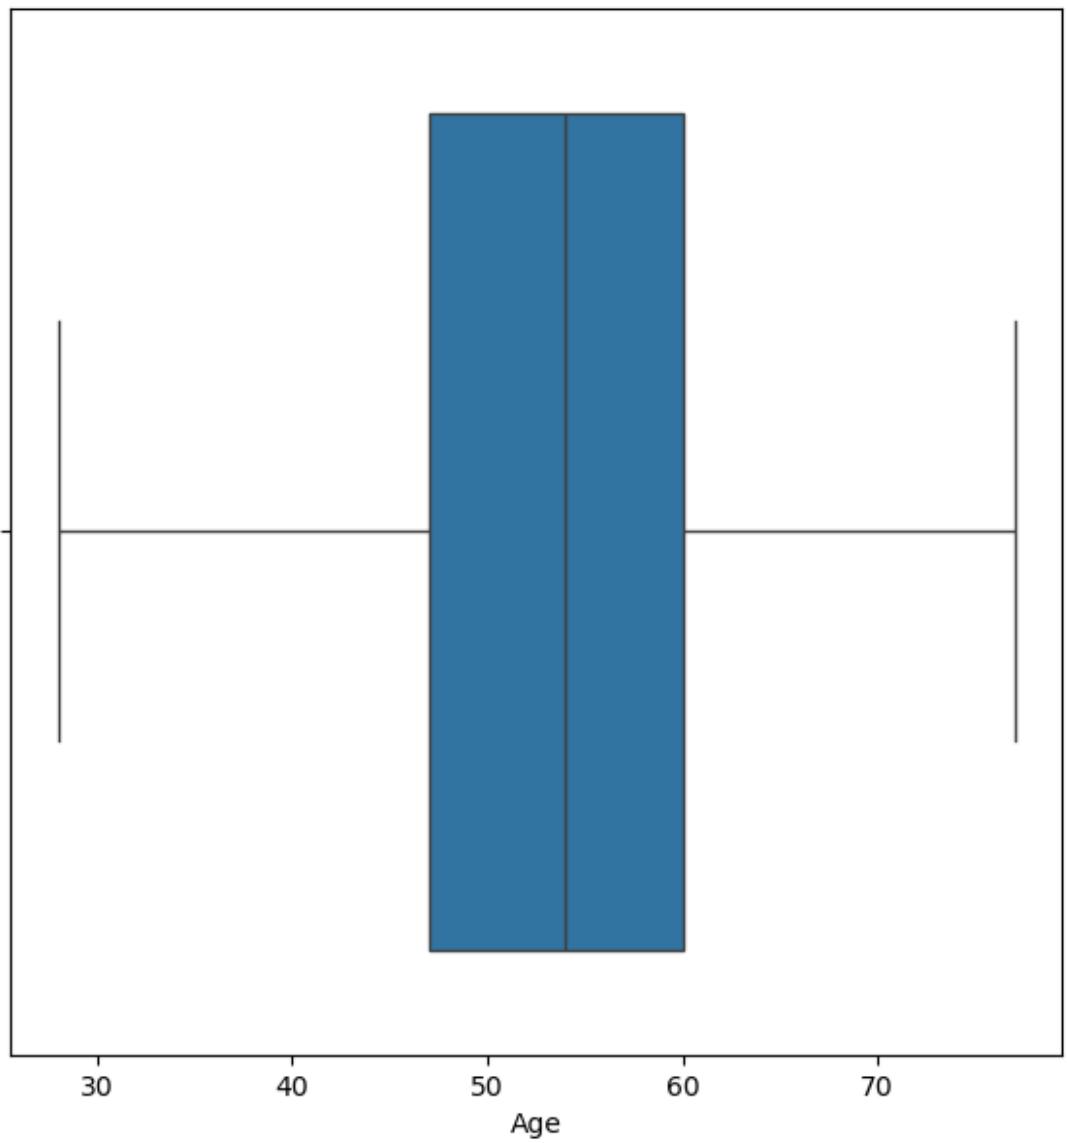
array(['M', 'F'], dtype=object)

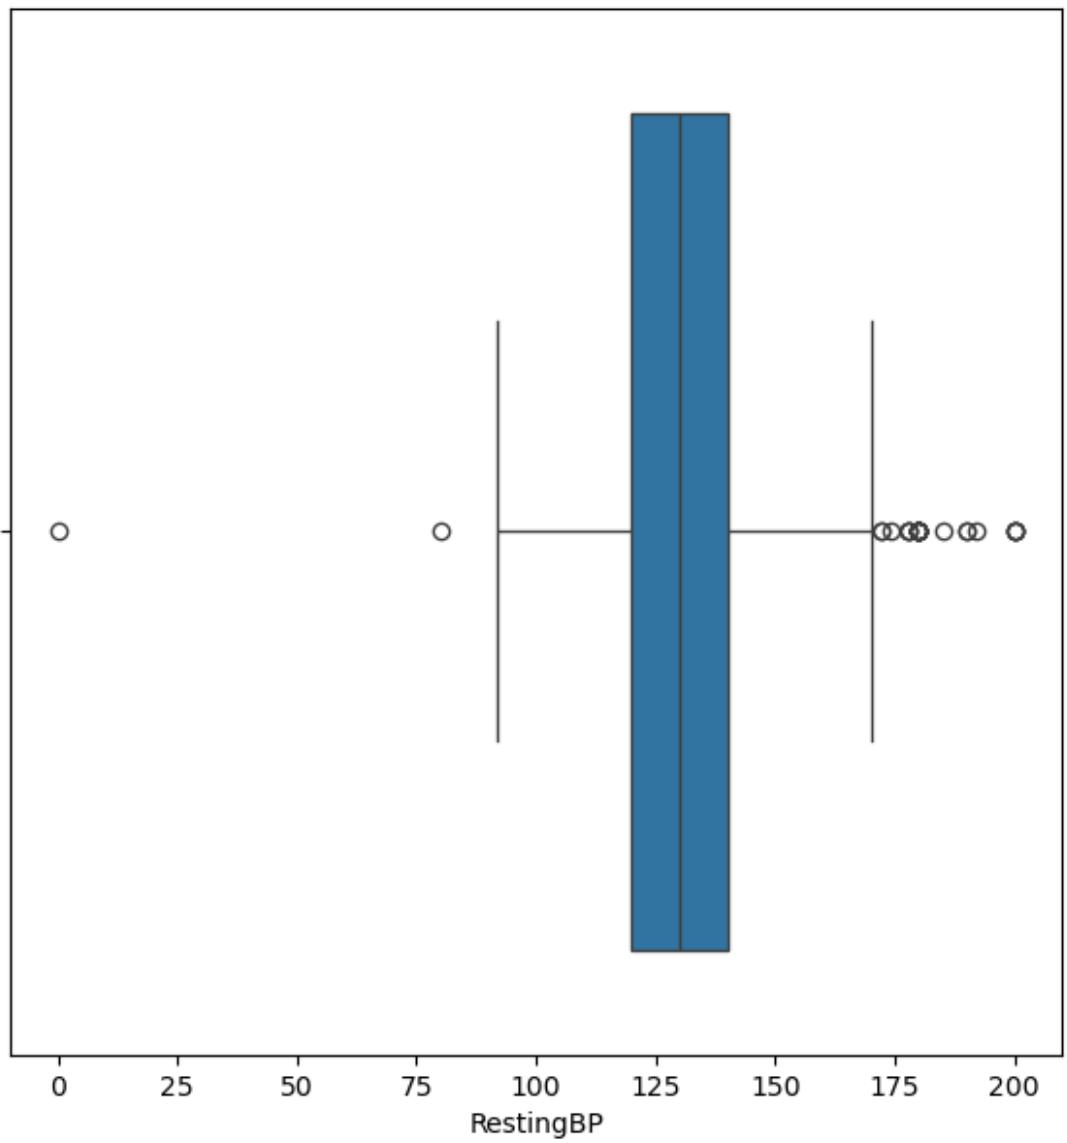
df['ChestPainType'].unique()

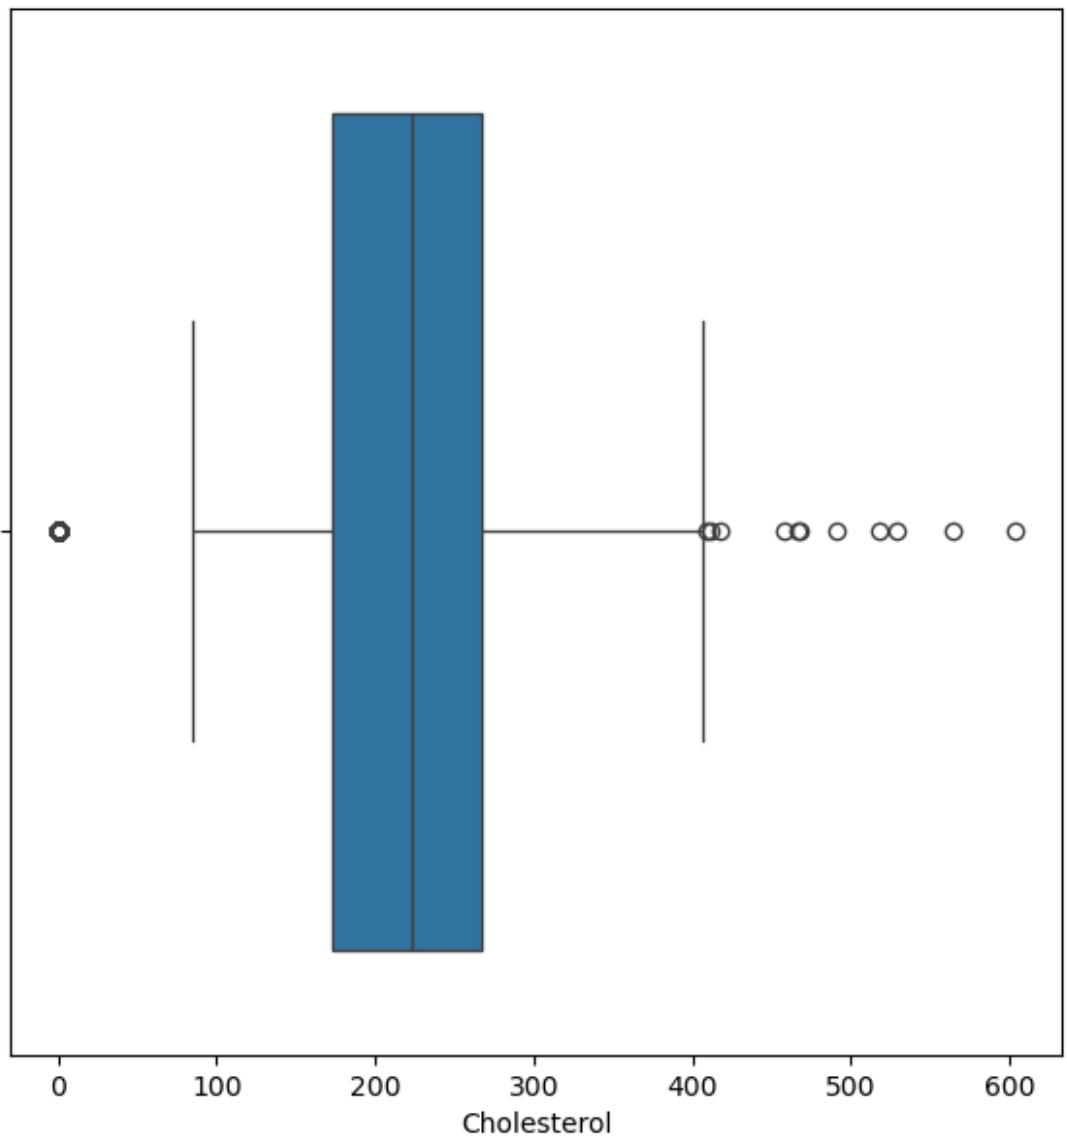
array(['ATA', 'NAP', 'ASY', 'TA'], dtype=object)

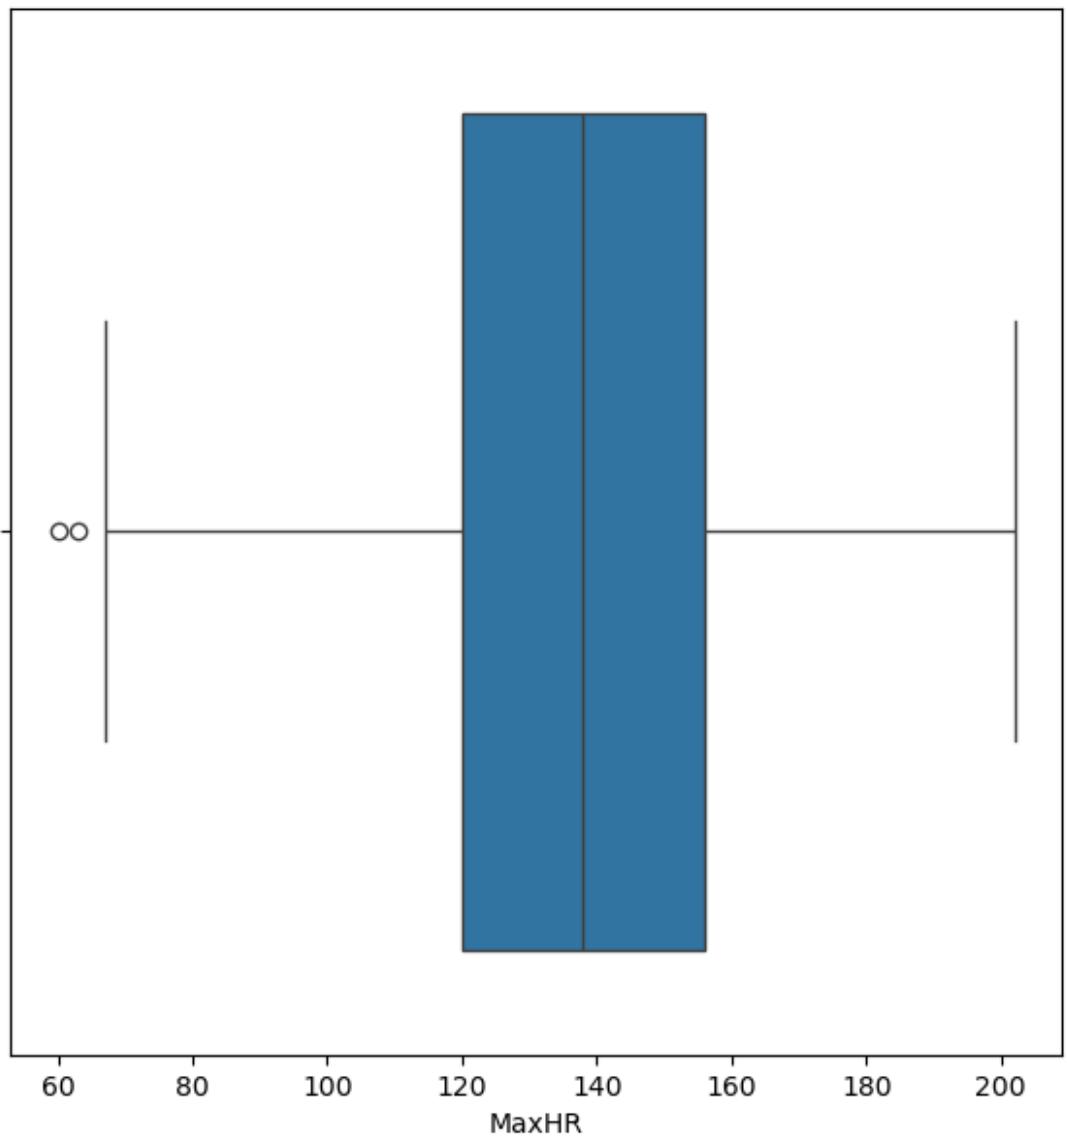
```

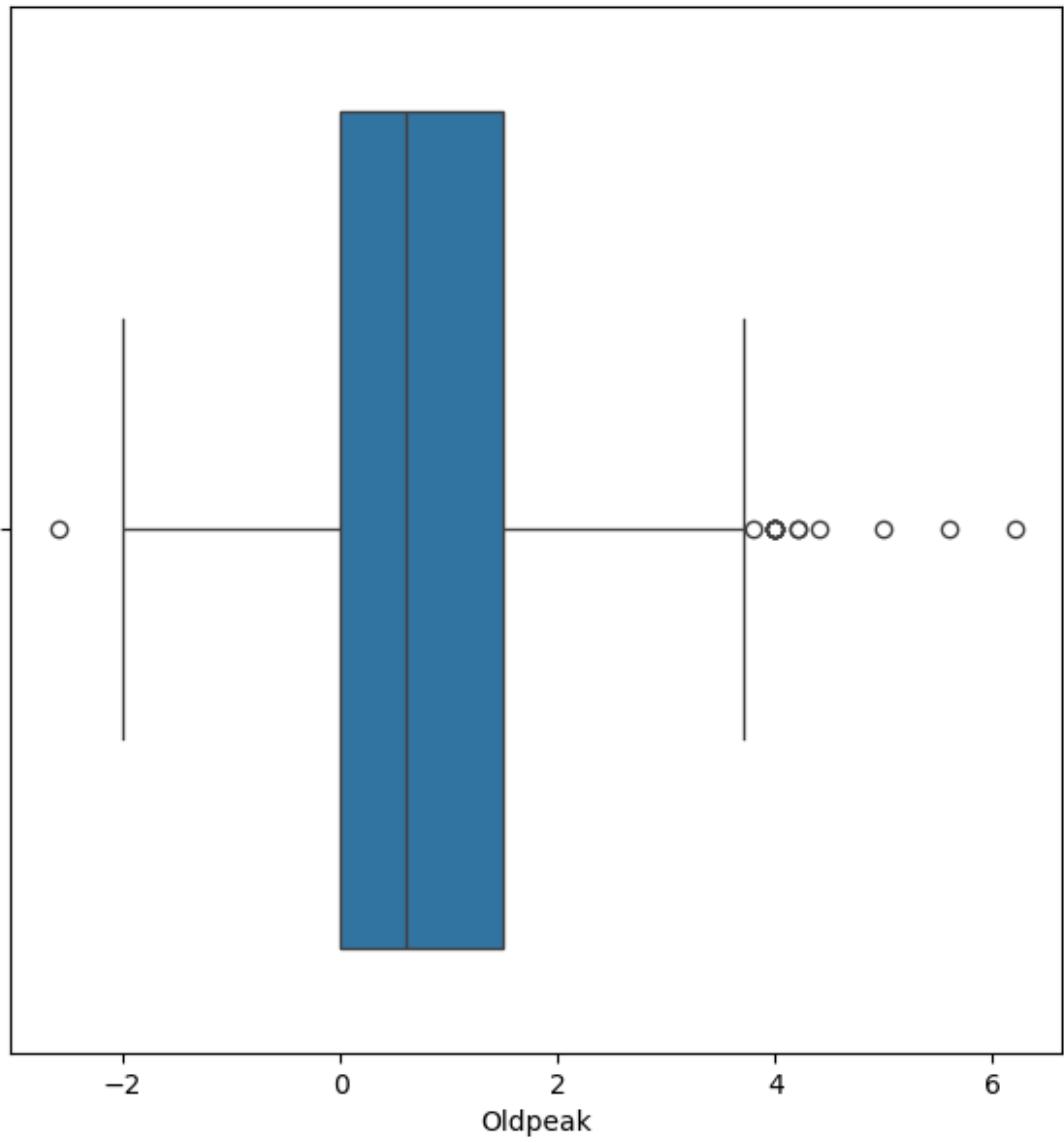
```
df['RestingECG'].unique()
array(['Normal', 'ST', 'LVH'], dtype=object)
df['ExerciseAngina'].unique()
array(['N', 'Y'], dtype=object)
df['ST_Slope'].unique()
array(['Up', 'Flat', 'Down'], dtype=object)
df['HeartDisease'].unique()
array([0, 1])
df['FastingBS'].unique()
array([0, 1])
df.duplicated().any()
np.False_
df.columns
Index(['Age', 'Sex', 'ChestPainType', 'RestingBP', 'Cholesterol',
       'FastingBS',
       'RestingECG', 'MaxHR', 'ExerciseAngina', 'Oldpeak', 'ST_Slope',
       'HeartDisease'],
      dtype='object')
cols=['Age', 'RestingBP', 'Cholesterol',
      'MaxHR', 'Oldpeak',
      ]
for i in cols:
    plt.figure(figsize=(7,7))
    sns.boxplot(x=df[i])
```



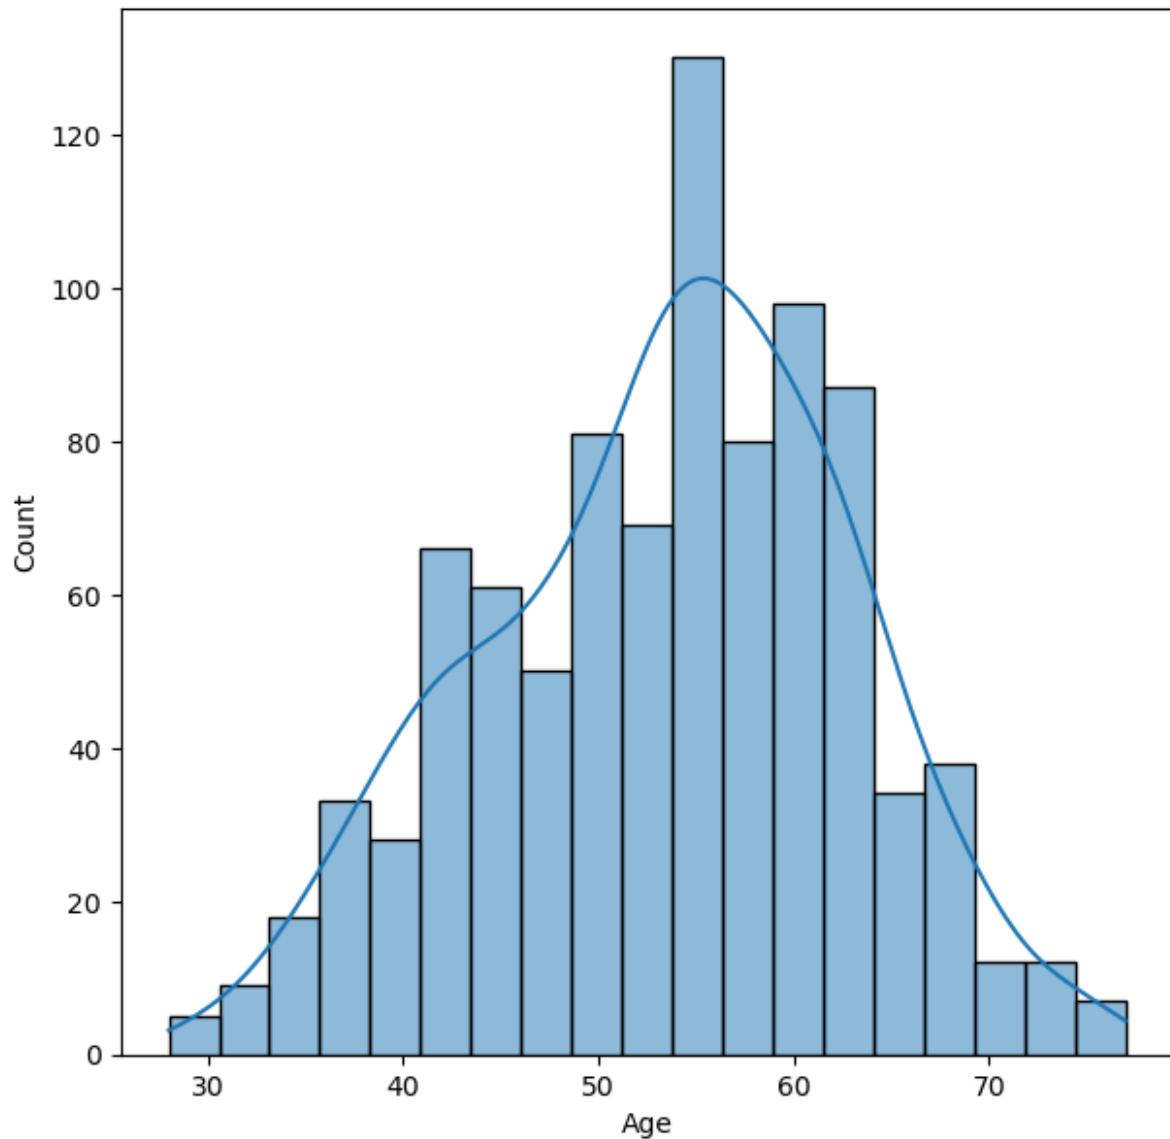


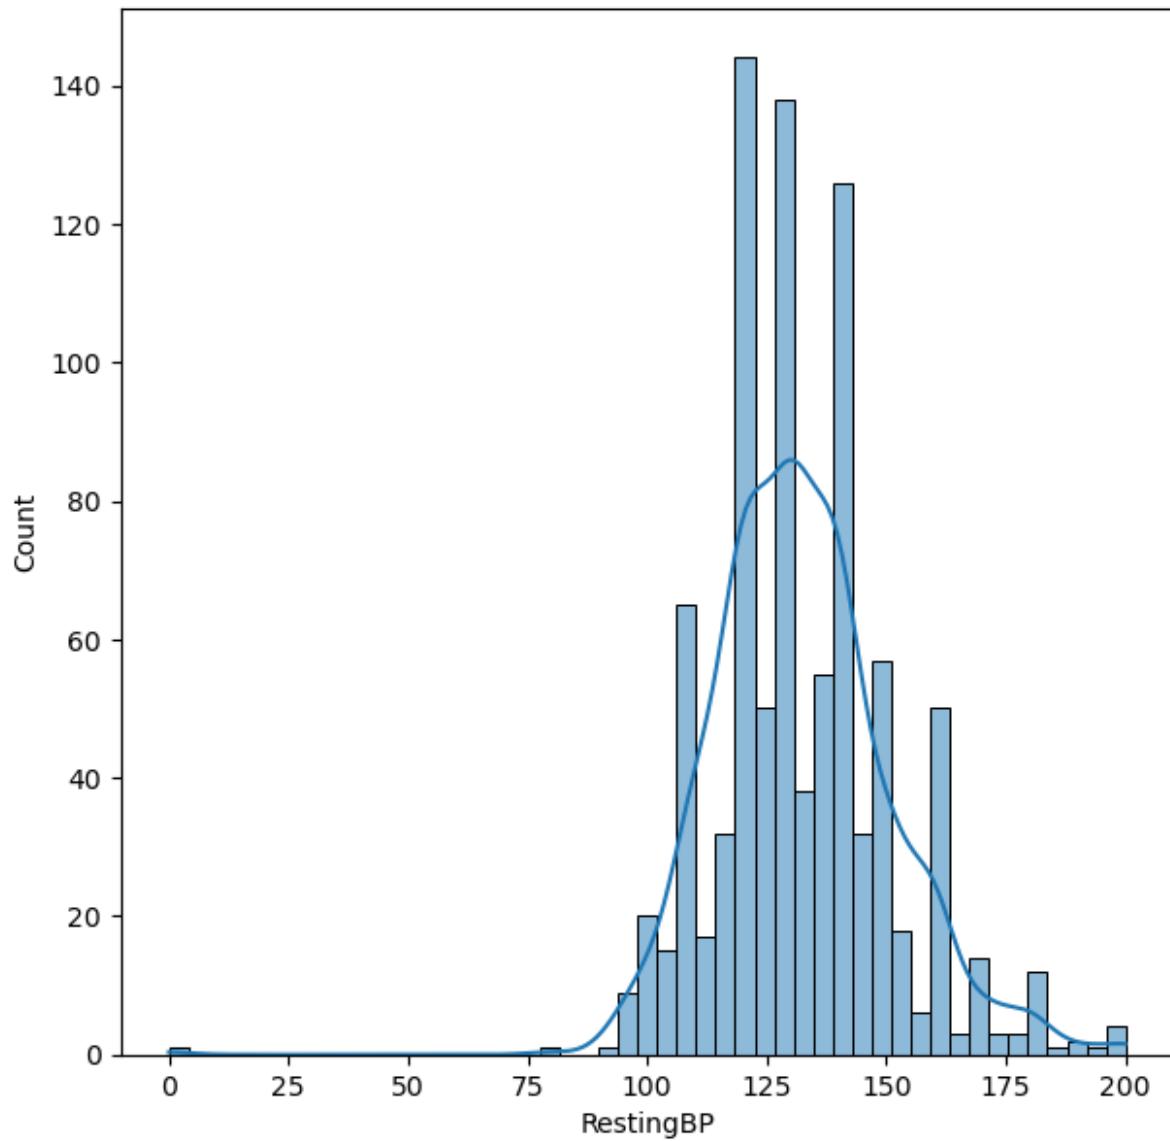


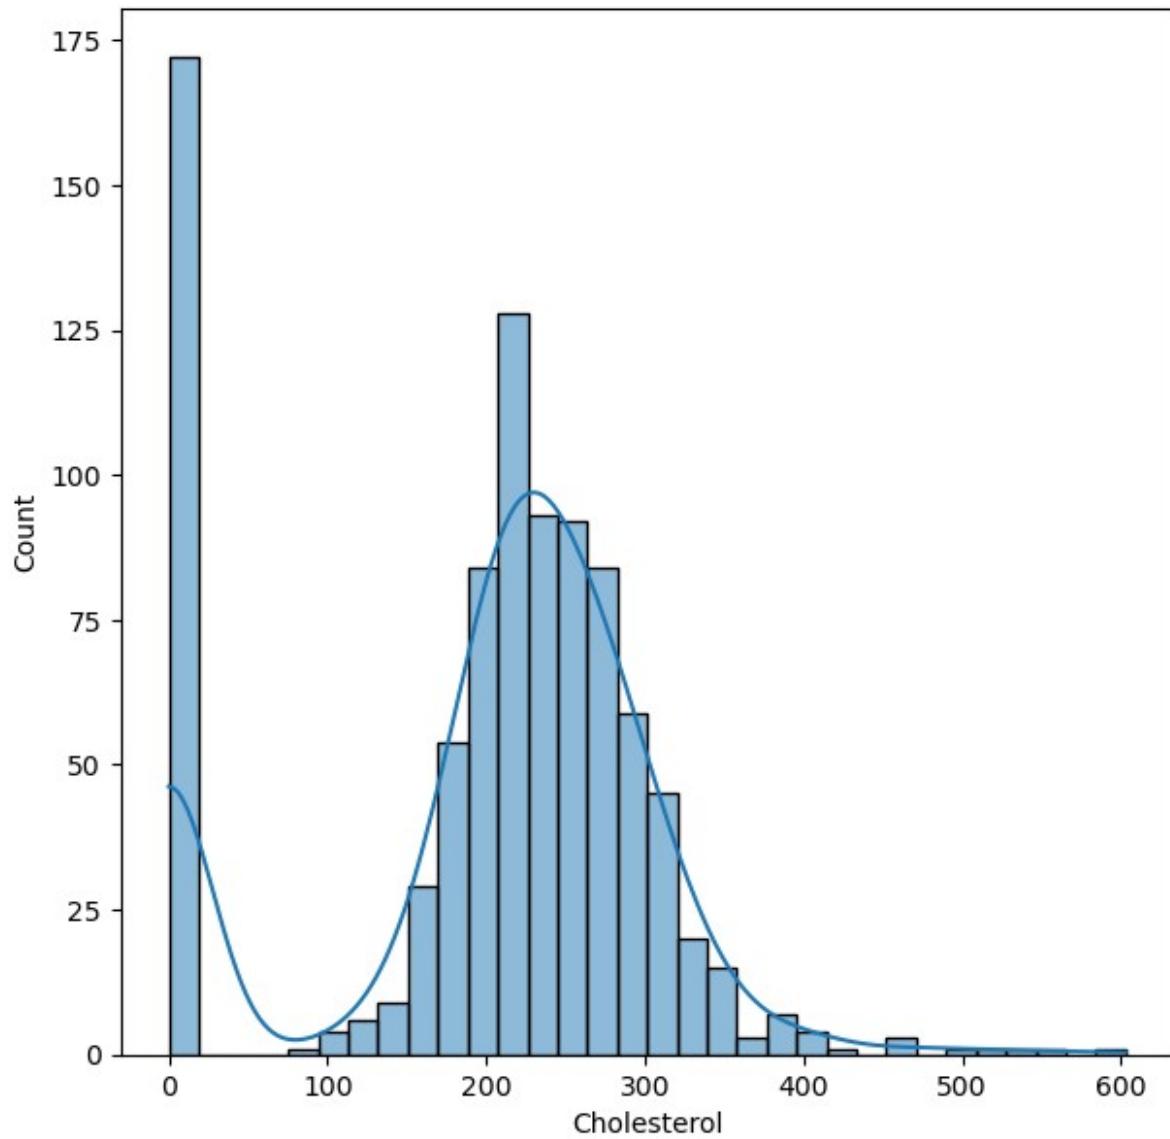


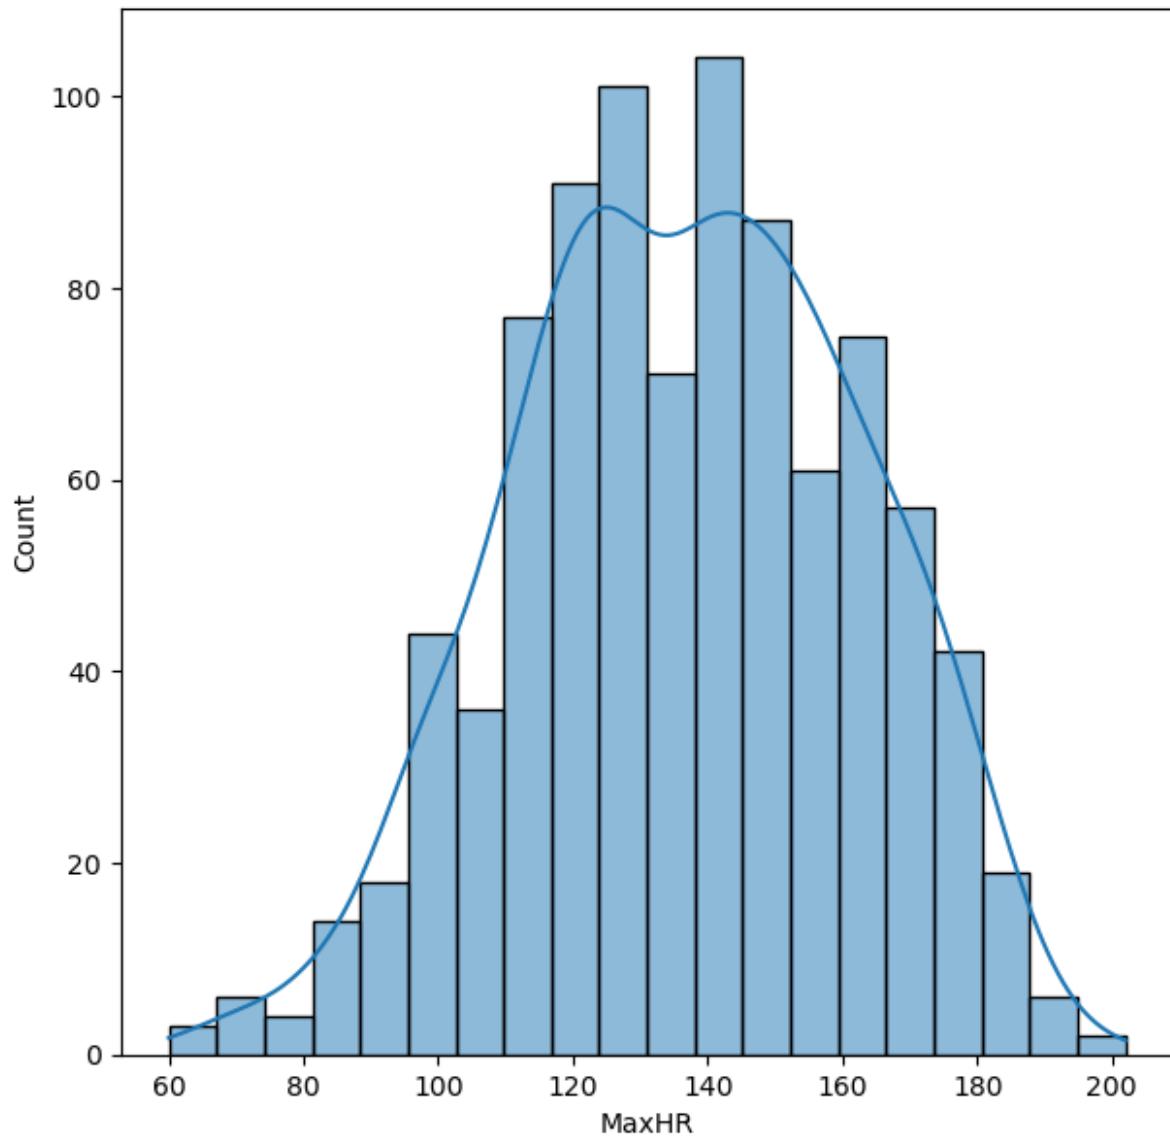


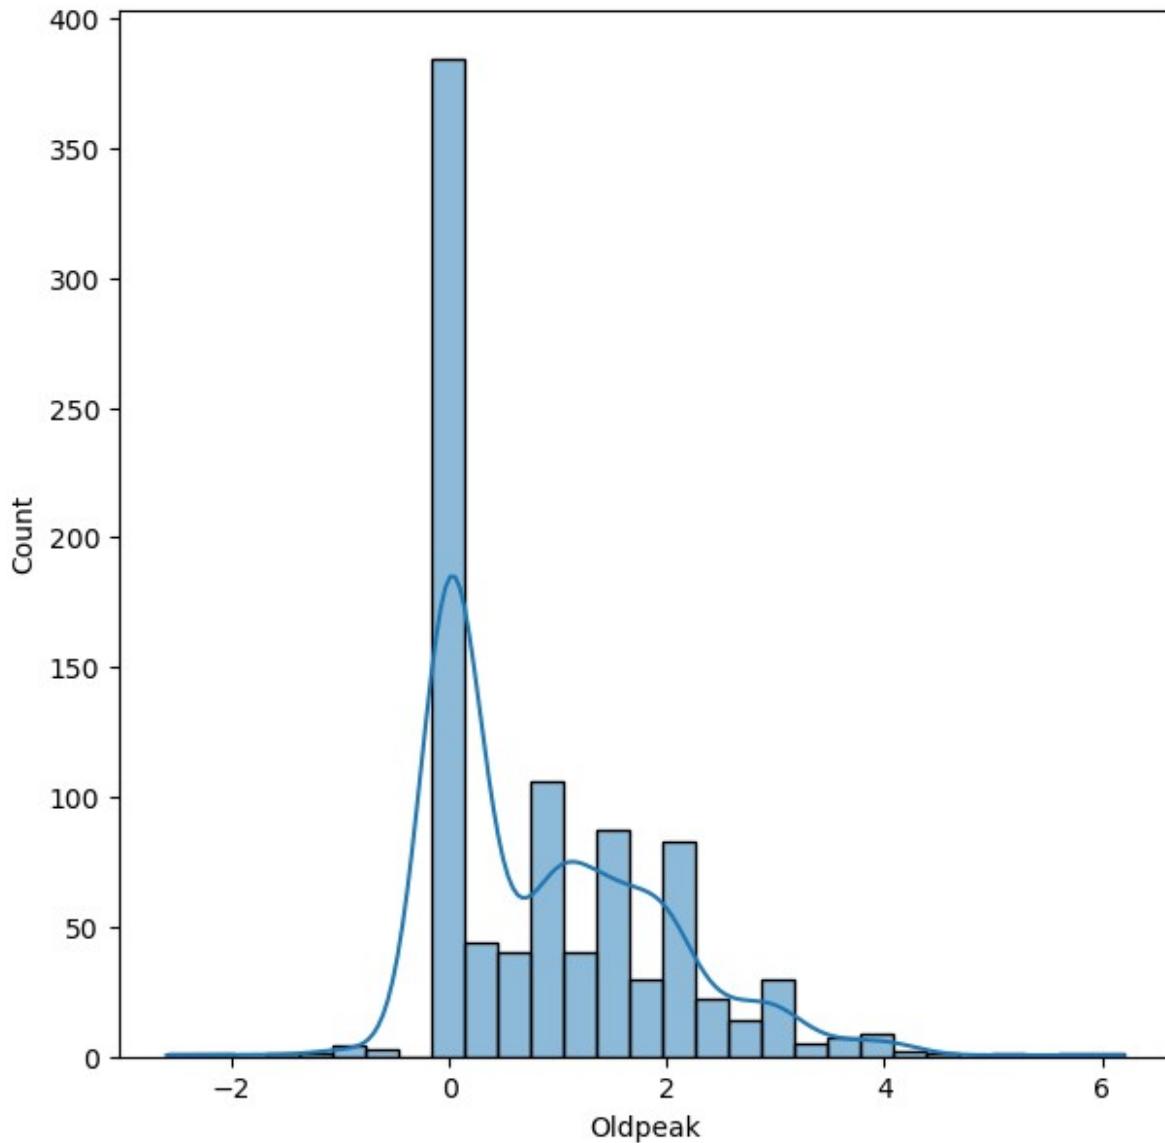
```
for i in cols:  
    plt.figure(figsize=(7,7))  
    sns.histplot(x=df[i],kde=True)
```



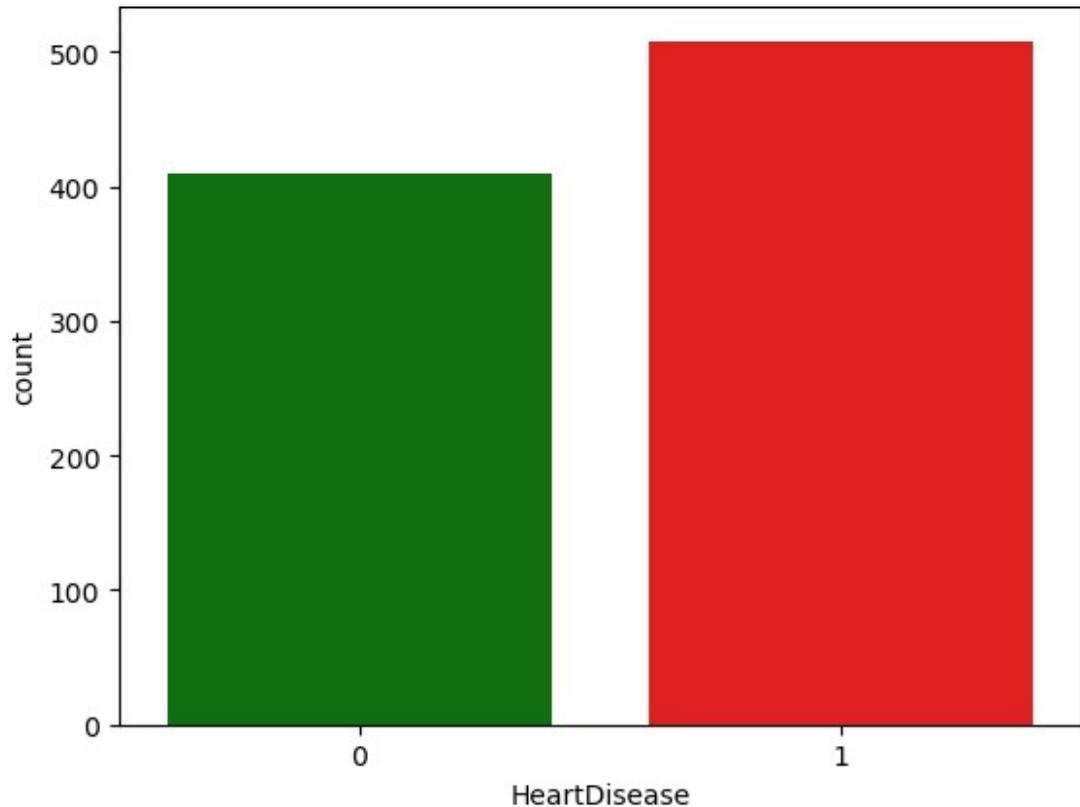






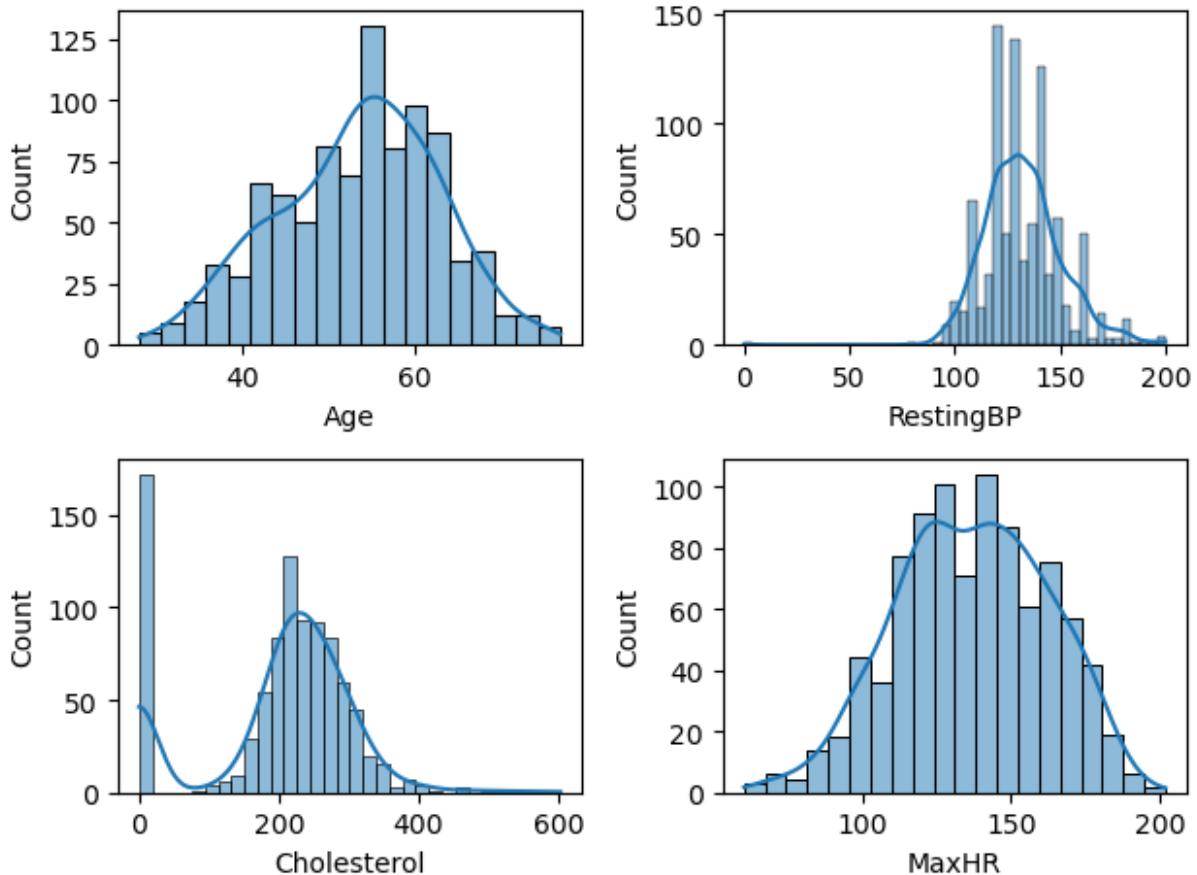


```
sns.countplot(x=df['HeartDisease'], palette=['green', 'red'])  
<Axes: xlabel='HeartDisease', ylabel='count'>
```



```
def plotting(var,num):
    plt.subplot(2,2,num)
    sns.histplot(df[var],kde=True)

plotting('Age',1)
plotting('RestingBP',2)
plotting('Cholesterol',3)
plotting('MaxHR',4)
plt.tight_layout()
```



```

df_clean=df.copy()

ch_mean=df_clean.loc[df_clean['Cholesterol'] != 0,'Cholesterol'].mean()

ch_mean

np.float64(244.6353887399464)

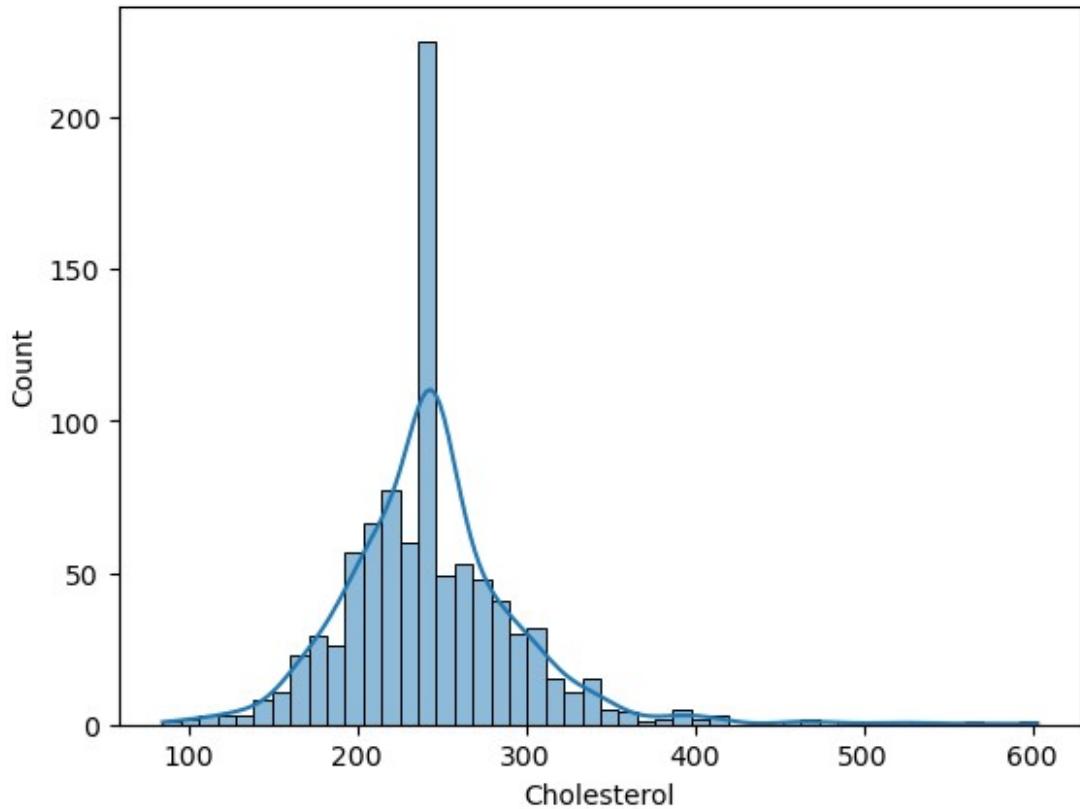
df_clean['Cholesterol']=df_clean['Cholesterol'].replace(0,ch_mean)
df_clean['Cholesterol']=df_clean['Cholesterol'].round(2)

ch_mean=df_clean.loc[df_clean['RestingBP'] != 0,'RestingBP'].mean()
df_clean['RestingBP']=df_clean['RestingBP'].replace(0,ch_mean)
df_clean['RestingBP']=df_clean['RestingBP'].round(2)

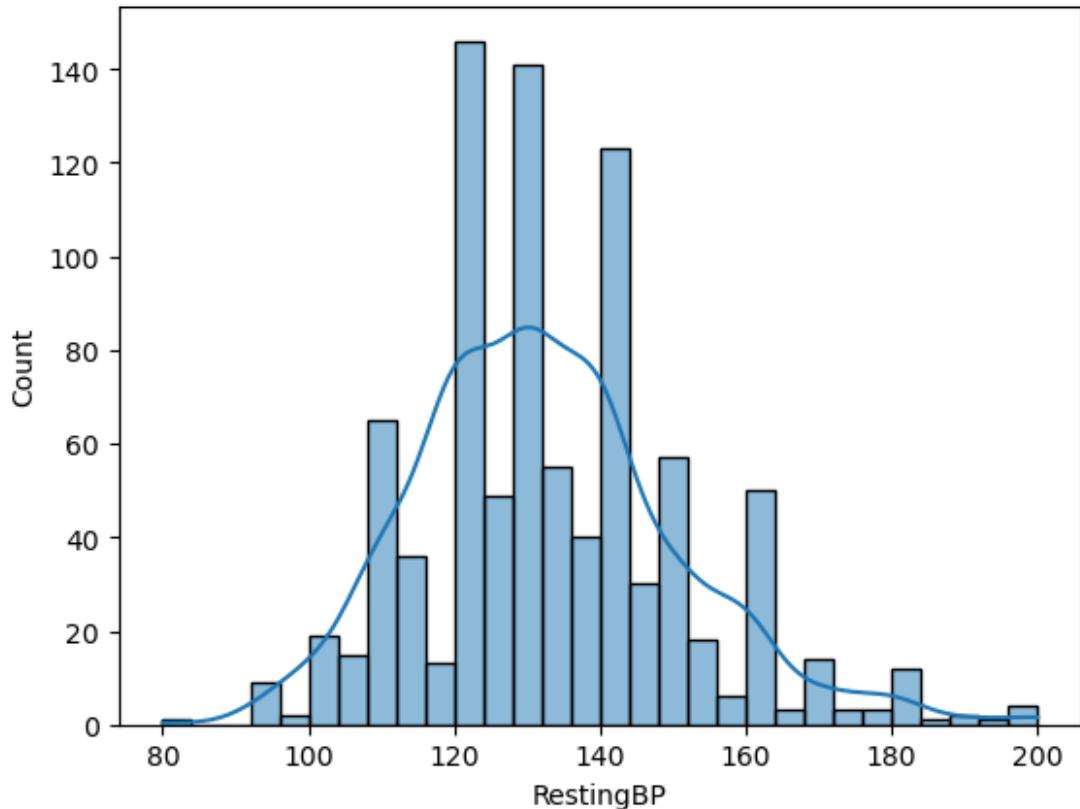
sns.histplot(df_clean['Cholesterol'],kde=True)

<Axes: xlabel='Cholesterol', ylabel='Count'>

```



```
sns.histplot(df_clean['RestingBP'], kde=True)  
<Axes: xlabel='RestingBP', ylabel='Count'>
```



```

df_clean['Sex']=df['Sex'].map({'M':1 , 'F':0})

df_clean.rename(columns={'Sex':'is_Male'},inplace=True)

df_clean

      Age  is_Male ChestPainType  RestingBP  Cholesterol  FastingBS
RestingECG \
0        40         1           ATA     140.0       289.0          0
Normal
1        49         0           NAP     160.0       180.0          0
Normal
2        37         1           ATA     130.0       283.0          0
ST
3        48         0           ASY     138.0       214.0          0
Normal
4        54         1           NAP     150.0       195.0          0
Normal
...
...
913     45         1           TA     110.0       264.0          0
Normal
914     68         1           ASY     144.0       193.0          1
Normal
915     57         1           ASY     130.0       131.0          0

```

Normal								
916	57	0	ATA	130.0	236.0	0		
LVH								
917	38	1	NAP	138.0	175.0	0		
Normal								
	MaxHR	ExerciseAngina	Oldpeak	ST_Slope	HeartDisease			
0	172	N	0.0	Up	0			
1	156	N	1.0	Flat	1			
2	98	N	0.0	Up	0			
3	108	Y	1.5	Flat	1			
4	122	N	0.0	Up	0			
..			
913	132	N	1.2	Flat	1			
914	141	N	3.4	Flat	1			
915	115	Y	1.2	Flat	1			
916	174	N	0.0	Flat	1			
917	173	N	0.0	Up	0			
[918 rows x 12 columns]								
df_clean=pd.get_dummies(df_clean,columns=['ChestPainType'],drop_first=False)								
df_clean								
	Age	is_Male	RestingBP	Cholesterol	FastingBS	RestingECG	MaxHR	
\								
0	40	1	140.0	289.0	0	Normal	172	
1	49	0	160.0	180.0	0	Normal	156	
2	37	1	130.0	283.0	0	ST	98	
3	48	0	138.0	214.0	0	Normal	108	
4	54	1	150.0	195.0	0	Normal	122	
..	
913	45	1	110.0	264.0	0	Normal	132	
914	68	1	144.0	193.0	1	Normal	141	
915	57	1	130.0	131.0	0	Normal	115	
916	57	0	130.0	236.0	0	LVH	174	
917	38	1	138.0	175.0	0	Normal	173	

	ExerciseAngina	Oldpeak	ST_Slope	HeartDisease	ChestPainType_ASY
0	N	0.0	Up	0	False
1	N	1.0	Flat	1	False
2	N	0.0	Up	0	False
3	Y	1.5	Flat	1	True
4	N	0.0	Up	0	False
..
913	N	1.2	Flat	1	False
914	N	3.4	Flat	1	True
915	Y	1.2	Flat	1	True
916	N	0.0	Flat	1	False
917	N	0.0	Up	0	False
	ChestPainType_ATA	ChestPainType_NAP	ChestPainType_TA		
0	True	False	False		
1	False	True	False		
2	True	False	False		
3	False	False	False		
4	False	True	False		
..		
913	False	False	True		
914	False	False	False		
915	False	False	False		
916	True	False	False		
917	False	True	False		

[918 rows x 15 columns]

```
df_clean=pd.get_dummies(df_clean,columns=['RestingECG'],drop_first=False)

df_clean=pd.get_dummies(df_clean,columns=['ExerciseAngina'],drop_first=False)

df_clean=pd.get_dummies(df_clean,columns=['ST_Slope'],drop_first=False)

df_clean
```

Oldpeak	Age	is_Male	RestingBP	Cholesterol	FastingBS	MaxHR
0	40	1	140.0	289.0	0	172
1	49	0	160.0	180.0	0	156
2	37	1	130.0	283.0	0	98
3	48	0	138.0	214.0	0	108
4	54	1	150.0	195.0	0	122
...
913	45	1	110.0	264.0	0	132
914	68	1	144.0	193.0	1	141
915	57	1	130.0	131.0	0	115
916	57	0	130.0	236.0	0	174
917	38	1	138.0	175.0	0	173
HeartDisease	ChestPainType_NAP	ChestPainType_ASY	ChestPainType_ATA			
False	0	False	True			
True	1	False	False			
False	0	False	True			
False	1	True	False			
True	0	False	False			
...			
913	1	False	False			
914	1	True	False			
915	1	True	False			
916	1	False	True			
917	0	False	False			
ChestPainType_TA	RestingECG_LVH	RestingECG_Normal				

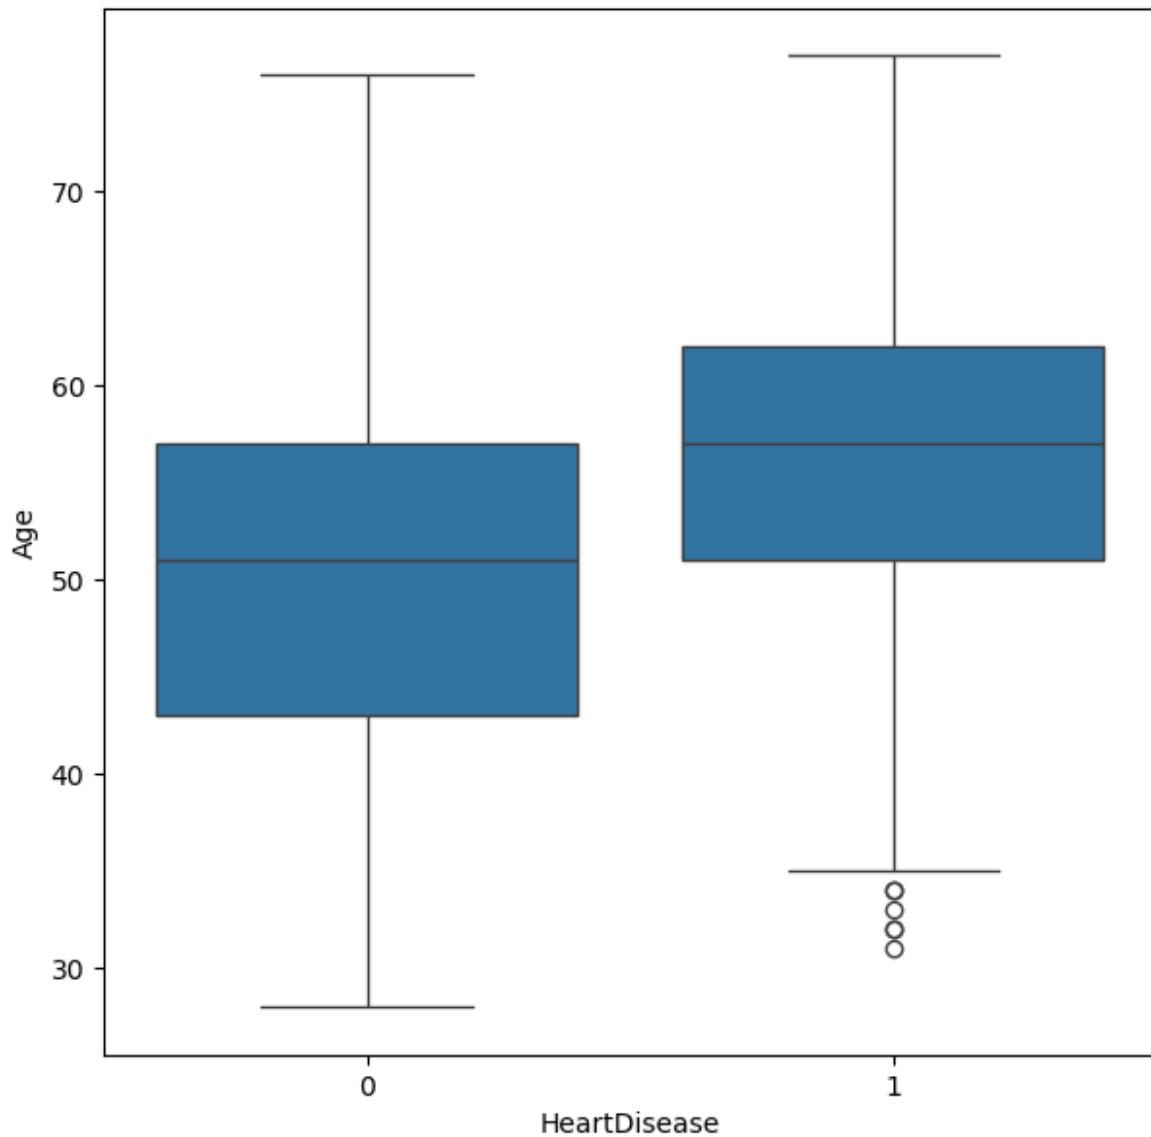
RestingECG_ST	\			
0	False	False	True	
1	False	False	True	
2	False	False	False	
3	False	False	True	
4	False	False	True	
...
.				
913	True	False	True	
914	False	False	True	
915	False	False	True	
916	False	True	False	
917	False	False	True	
False				
ExerciseAngina_N	ExerciseAngina_Y	ST_Slope_Down	ST_Slope_Flat	
\				
0	True	False	False	False
1	True	False	False	True
2	True	False	False	False
3	False	True	False	True
4	True	False	False	False
...
913	True	False	False	True
914	True	False	False	True
915	False	True	False	True
916	True	False	False	True
917	True	False	False	False
ST_Slope_Up				
0	True			

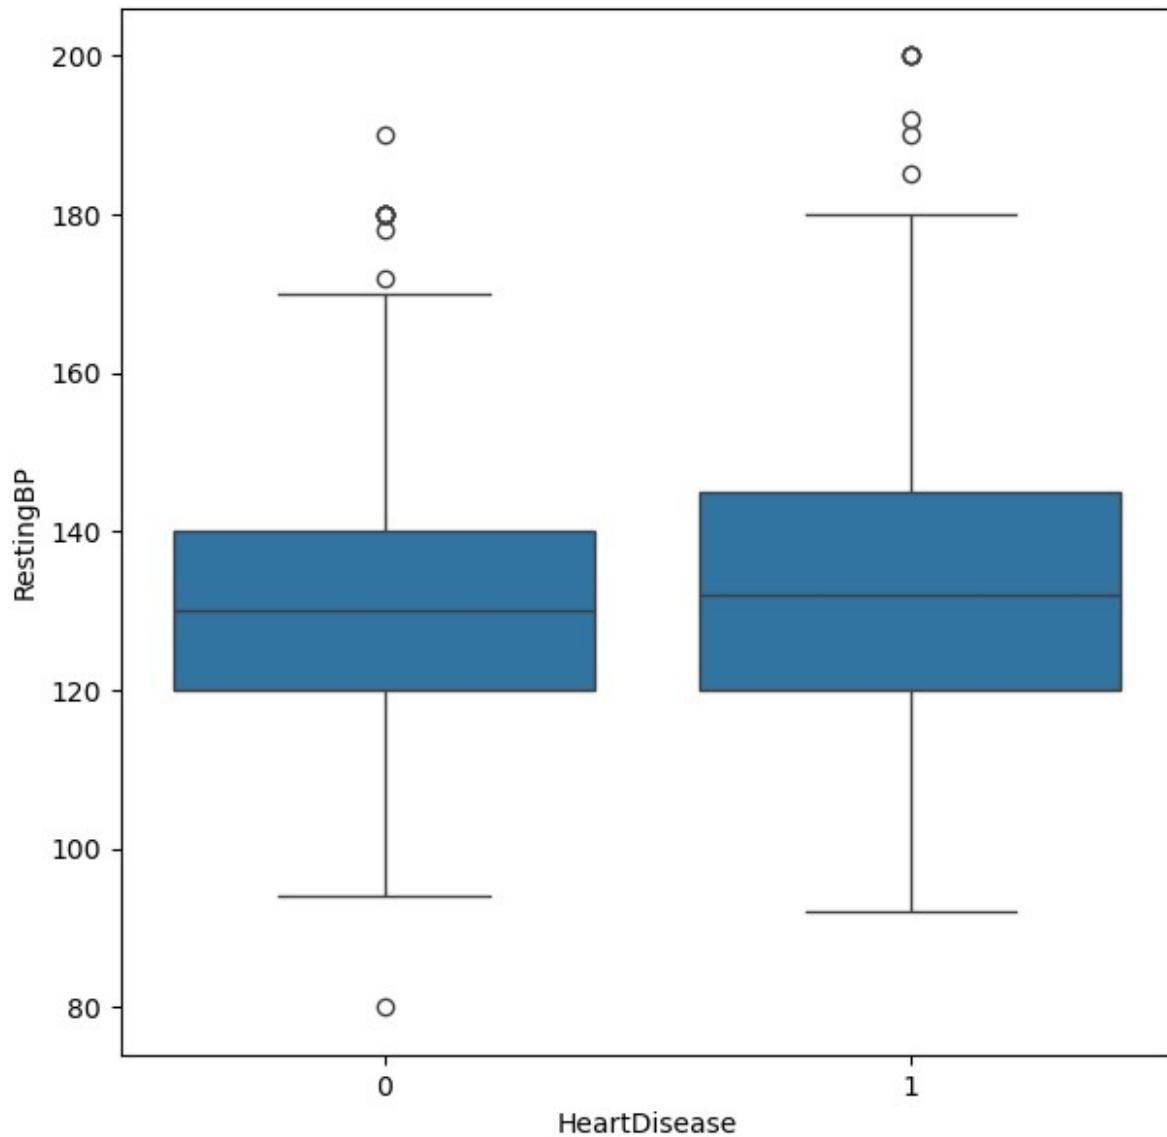
```
1      False
2      True
3      False
4      True
..
913     ...
914     False
915     False
916     False
917     True

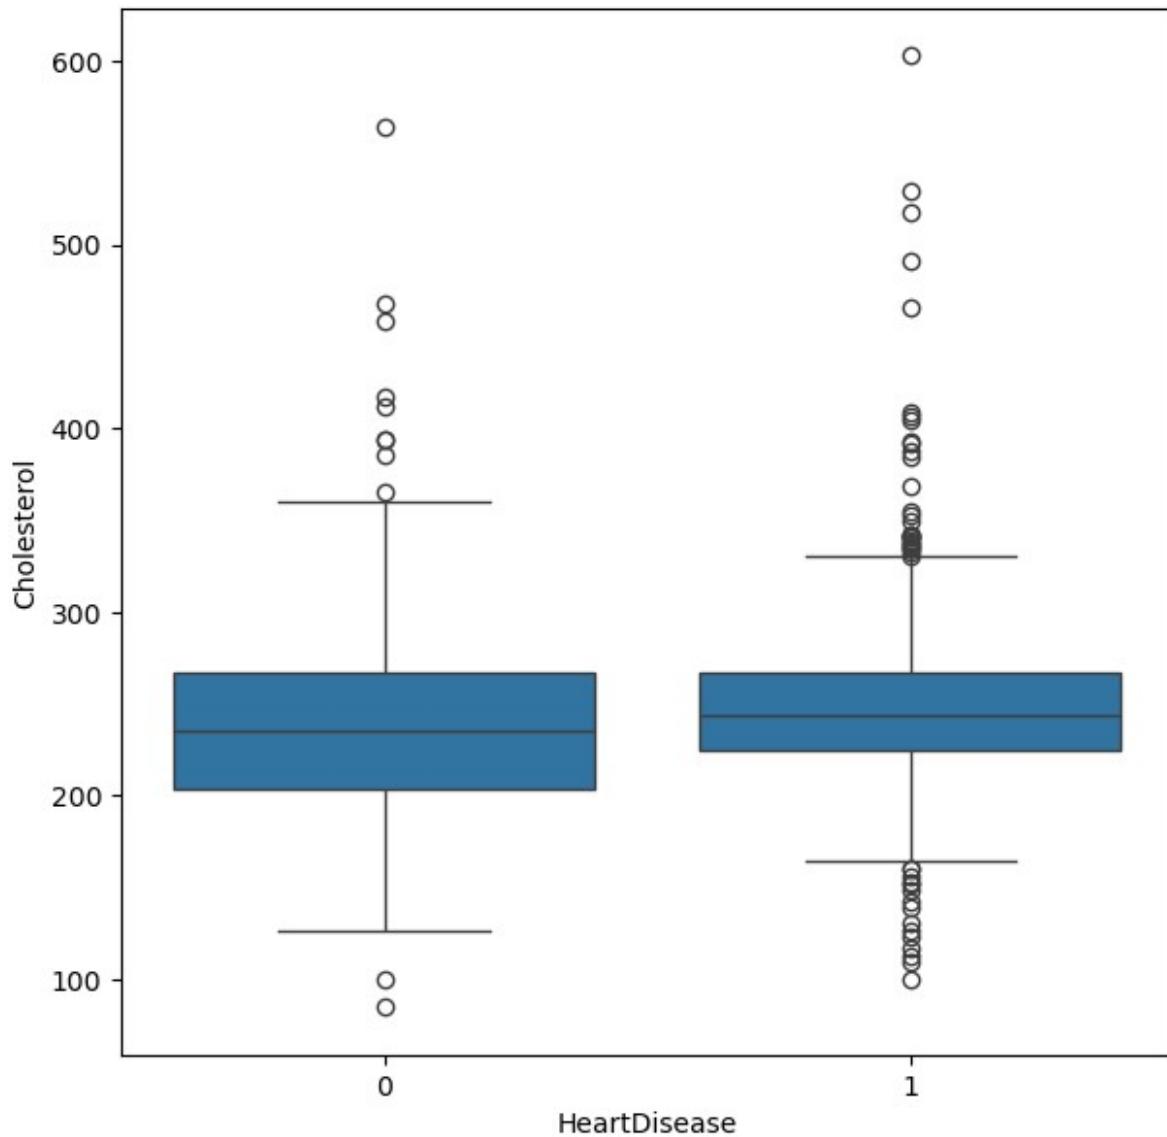
[918 rows x 20 columns]

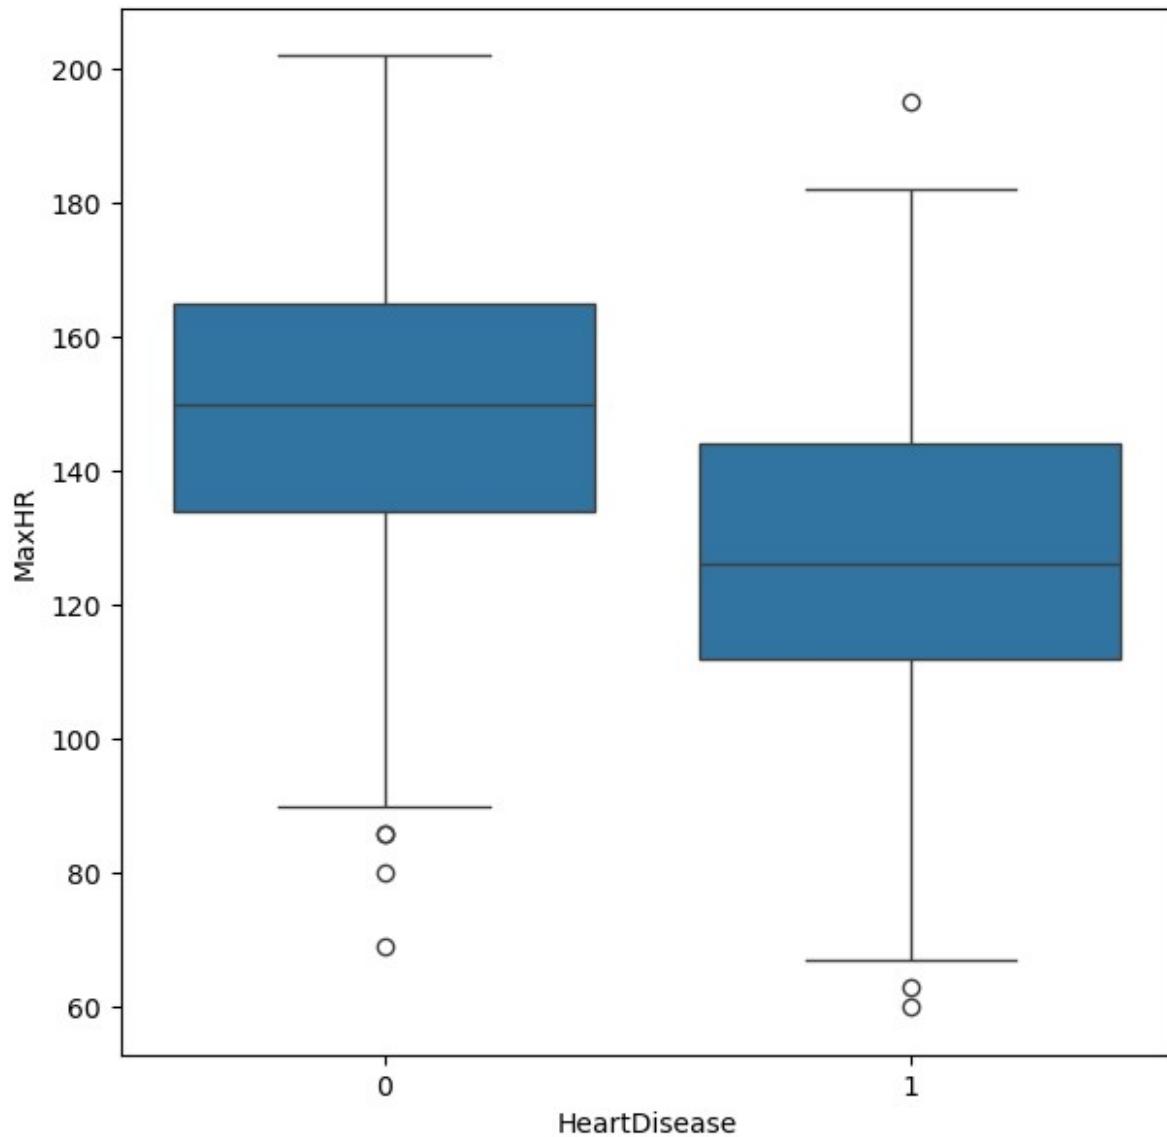
df_clean=df_clean.astype(int)

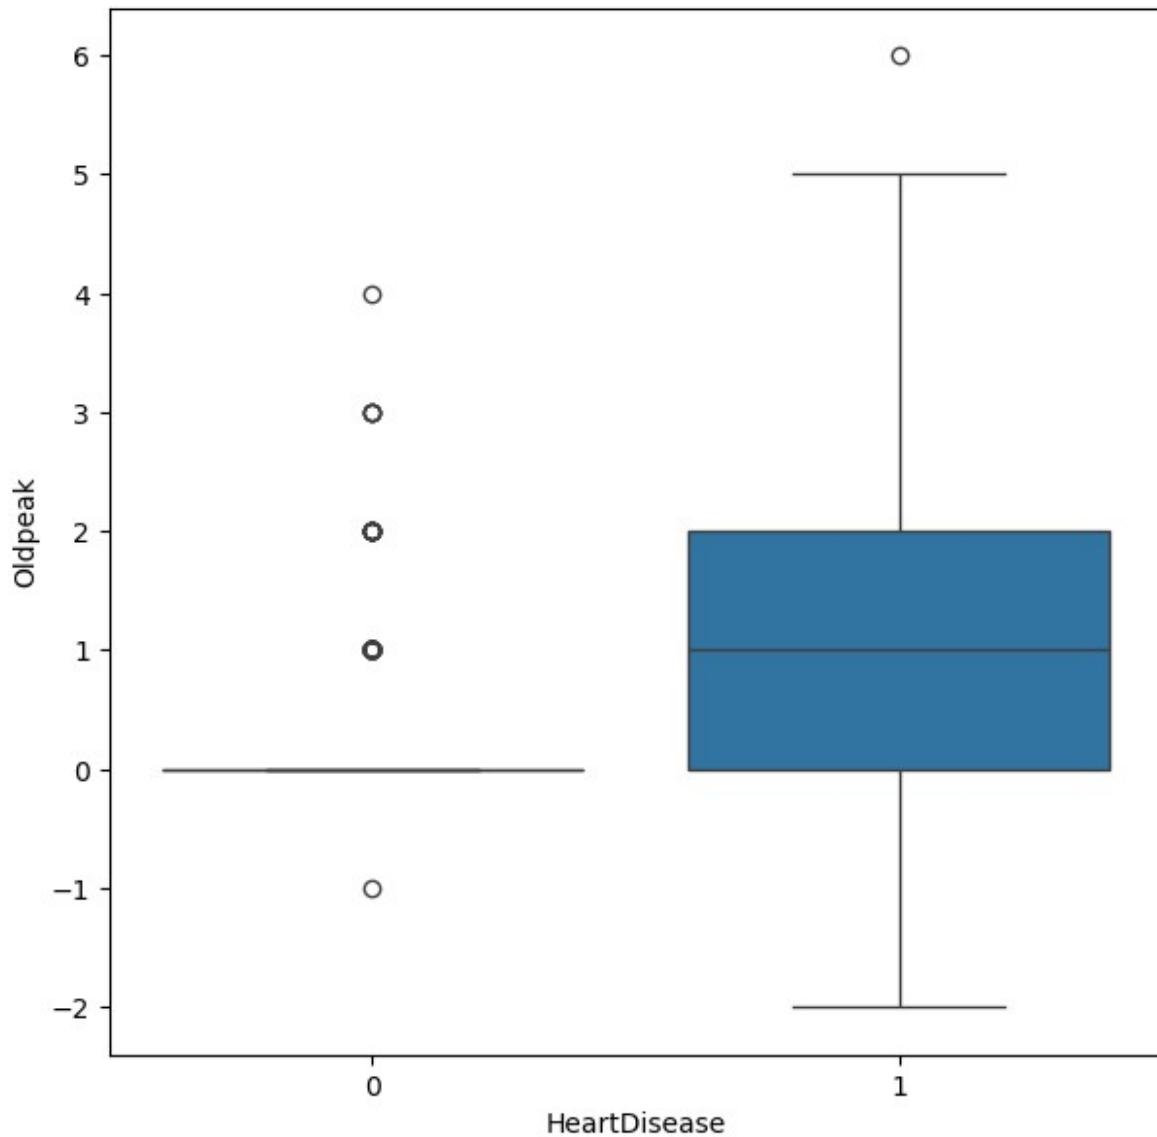
numeric_cols=['Age', 'RestingBP', 'Cholesterol', 'MaxHR', 'Oldpeak',
'HeartDisease']
for i in numeric_cols:
    plt.figure(figsize=(7,7))
    sns.boxplot(y=i,x='HeartDisease',data=df_clean)
```

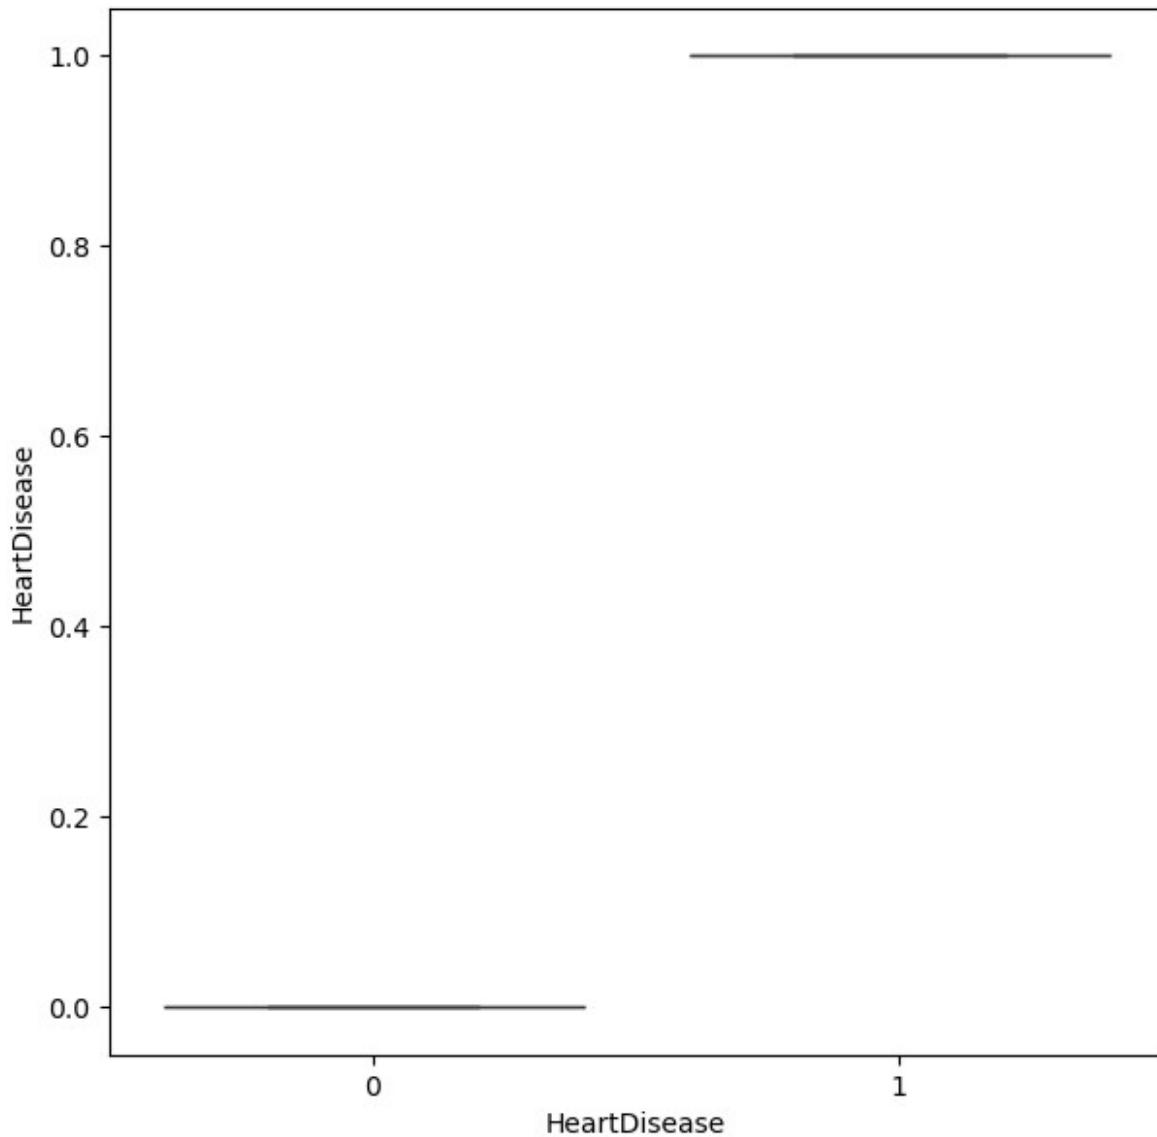




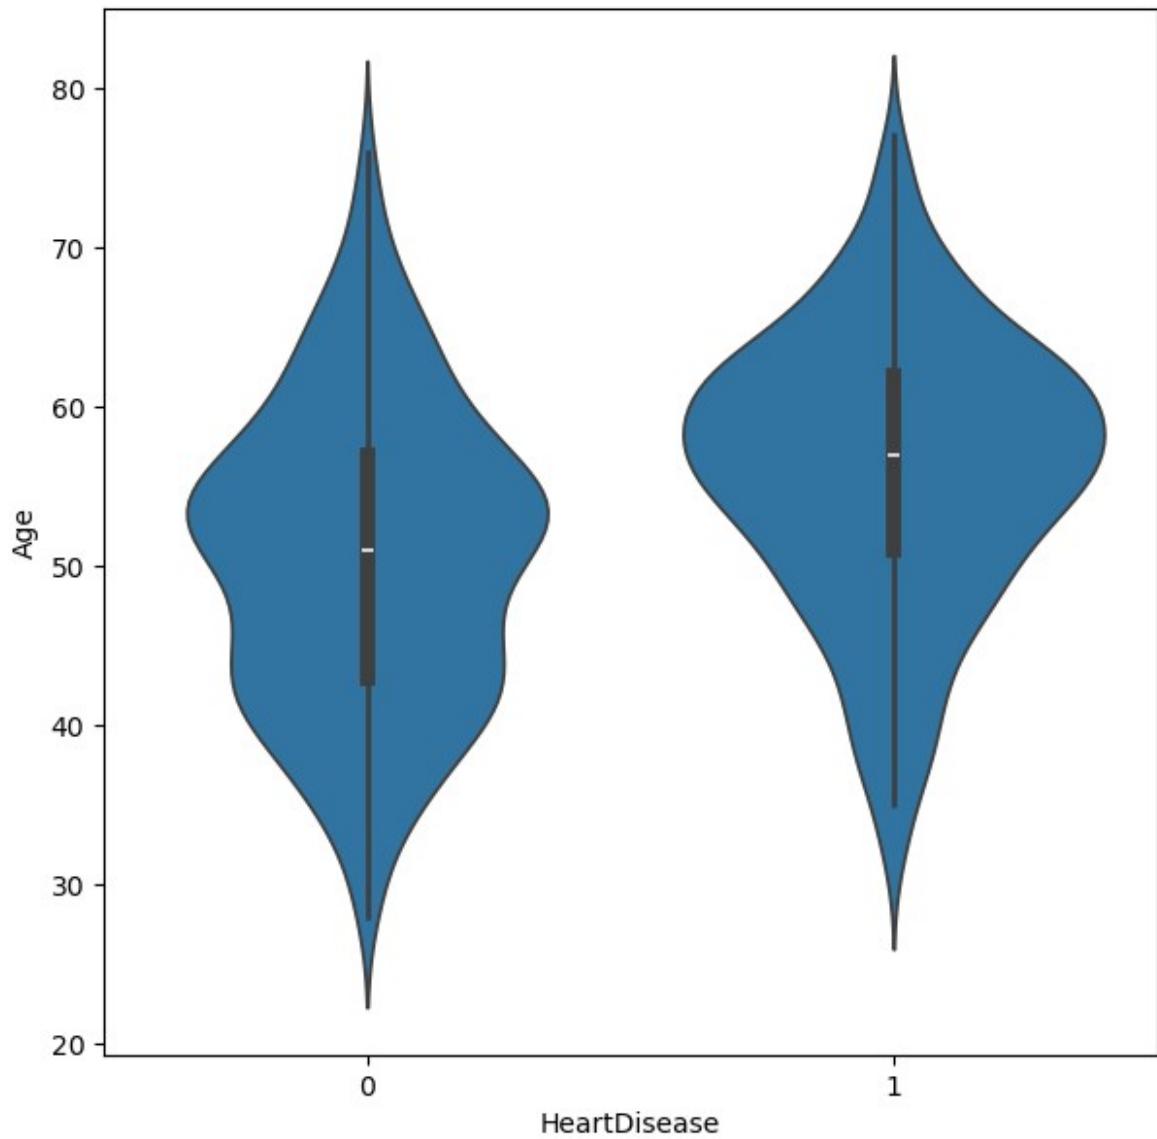


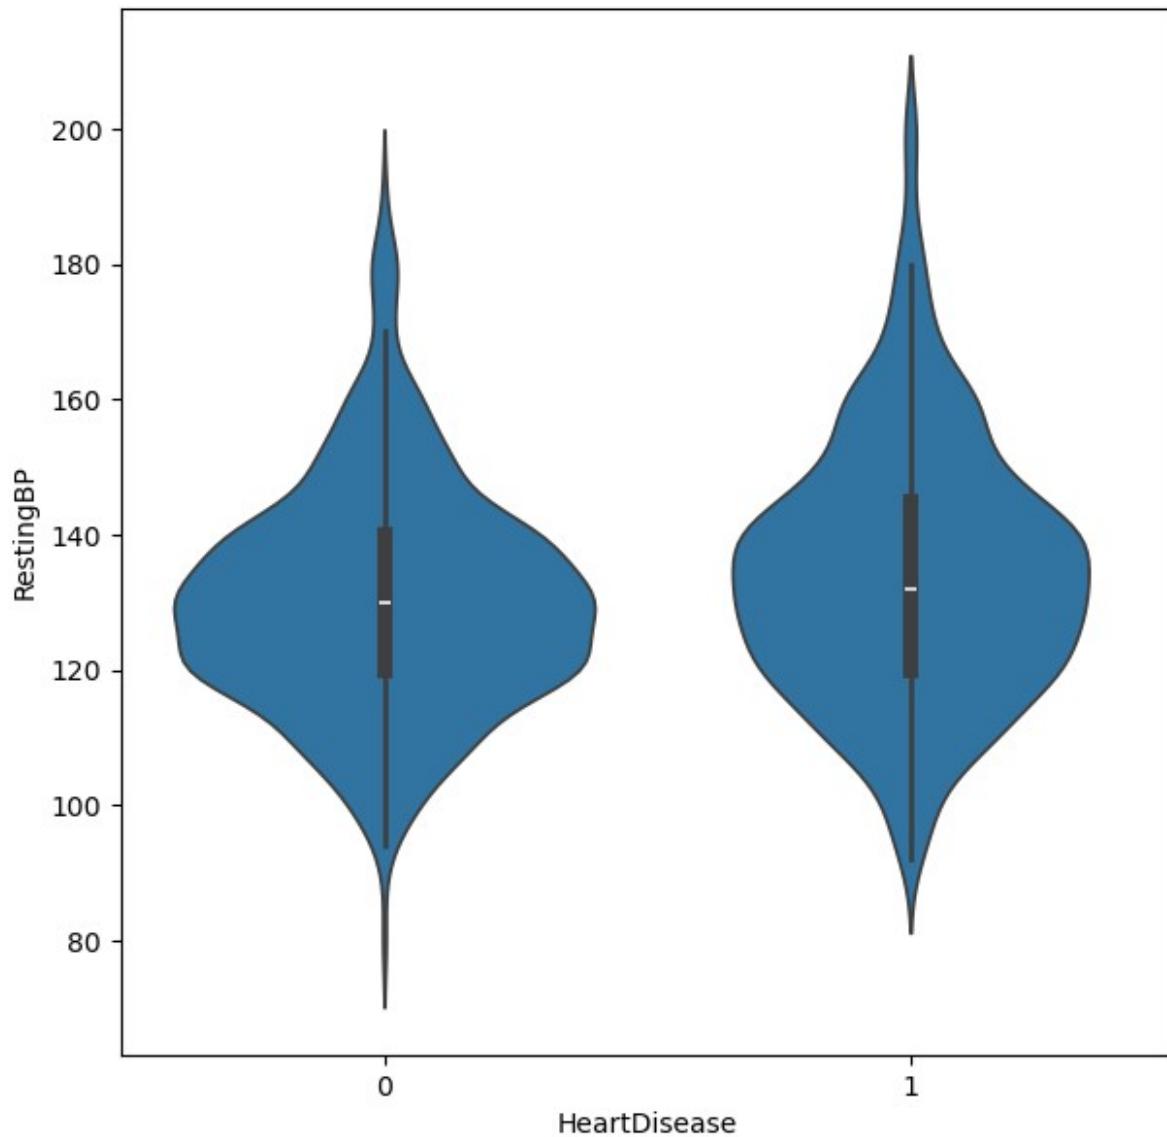


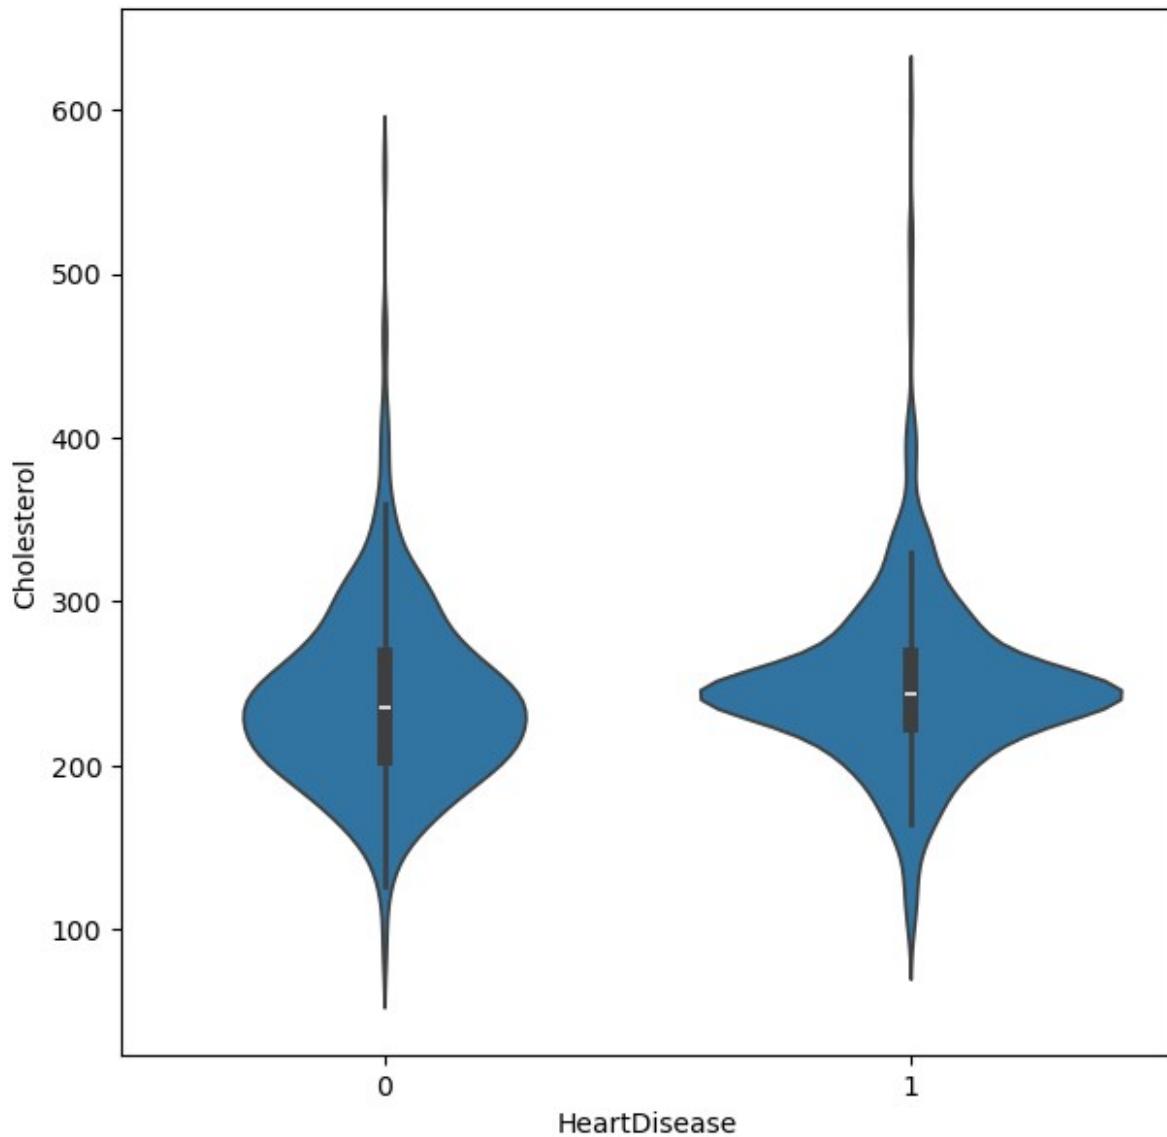


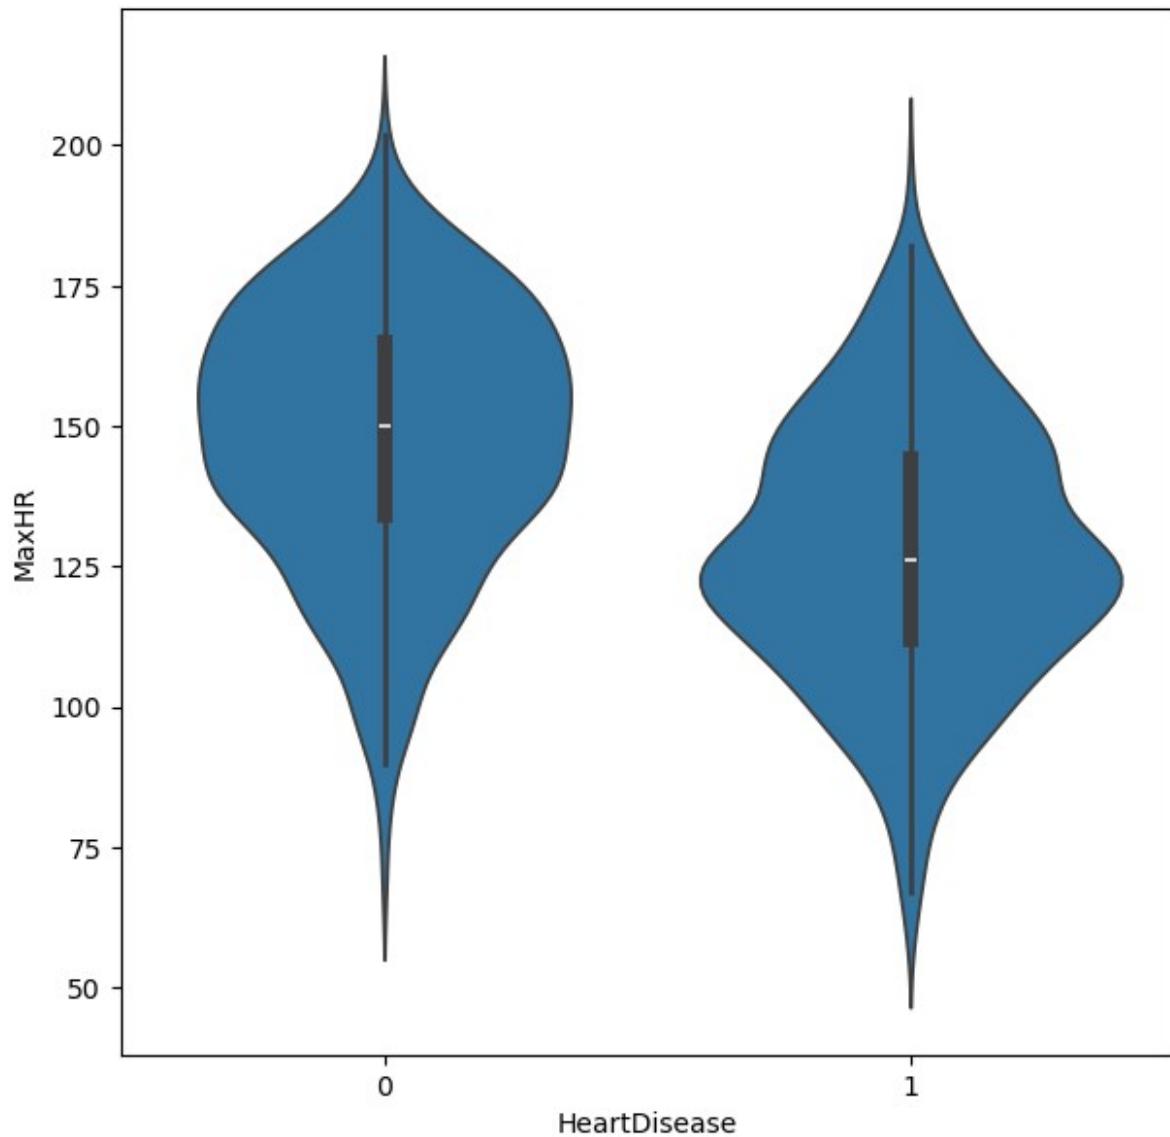


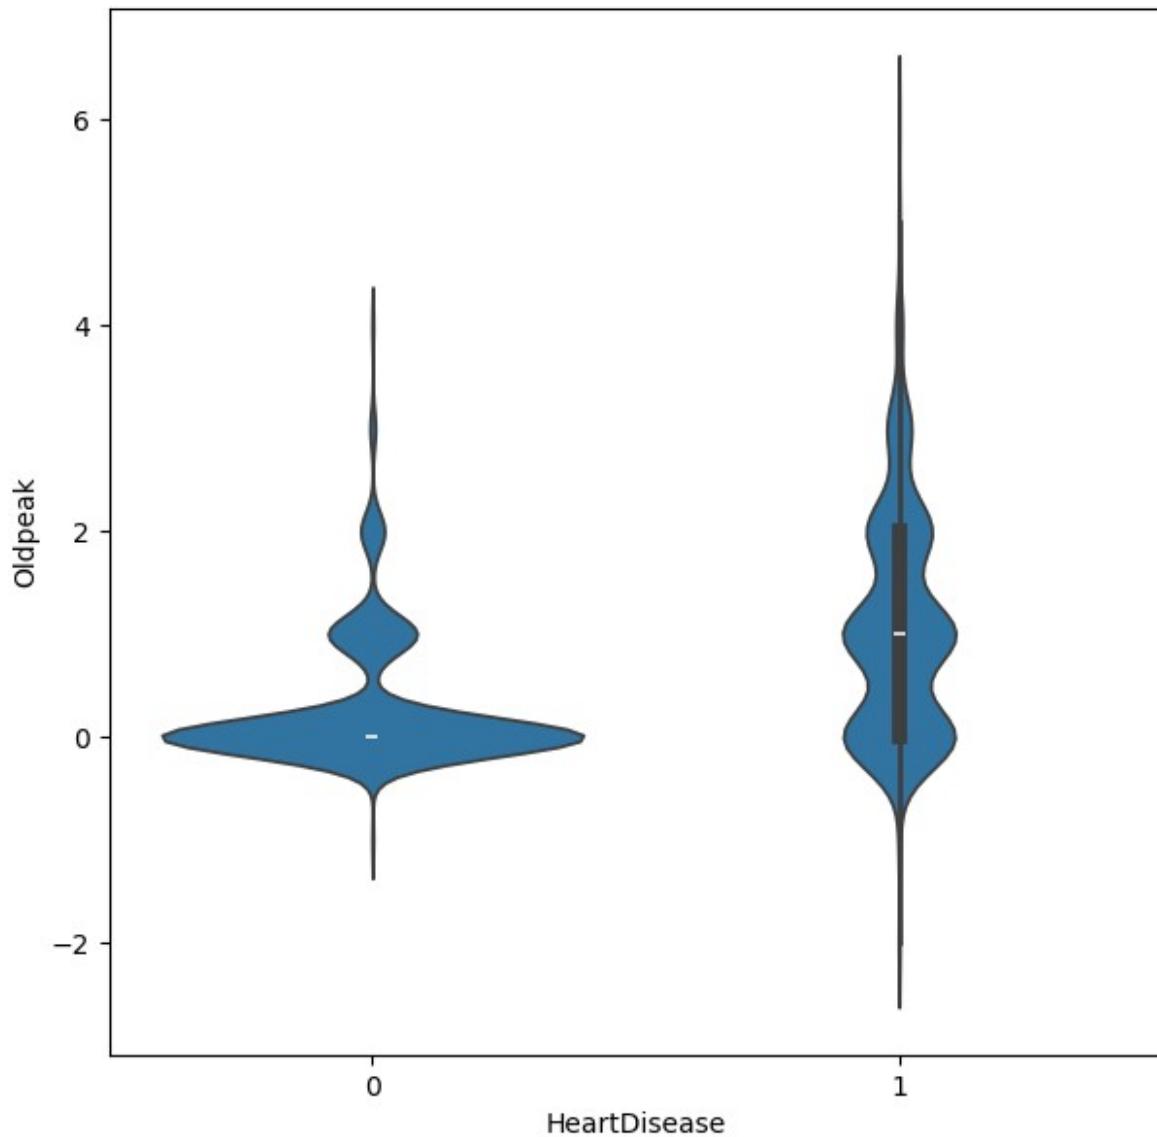
```
numeric_cols=['Age', 'RestingBP', 'Cholesterol', 'MaxHR', 'Oldpeak',
'HeartDisease']
for i in numeric_cols:
    plt.figure(figsize=(7,7))
    sns.violinplot(y=i,x='HeartDisease',data=df_clean)
```

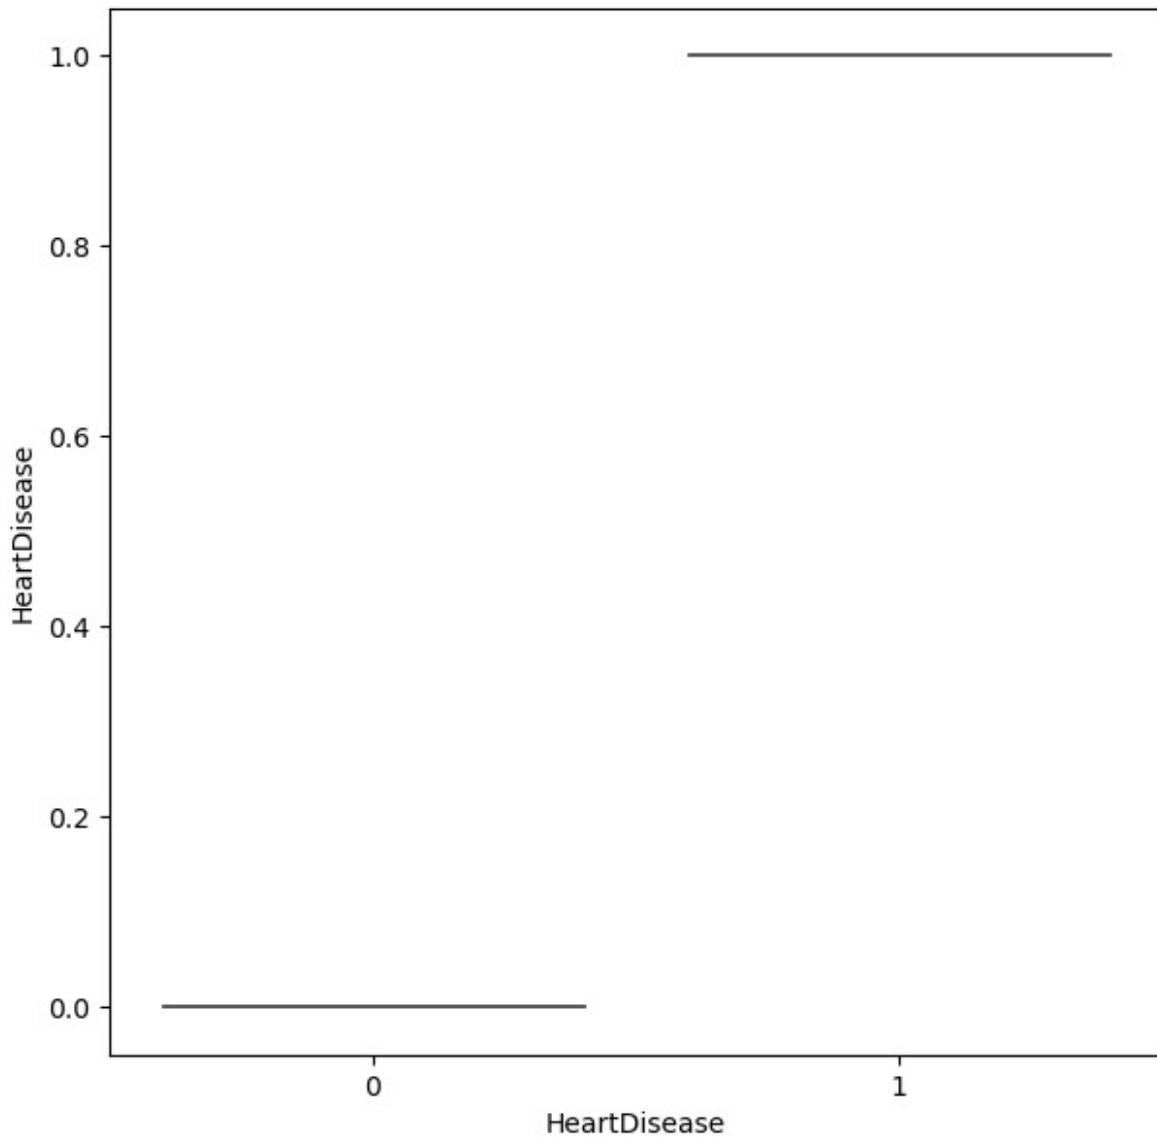








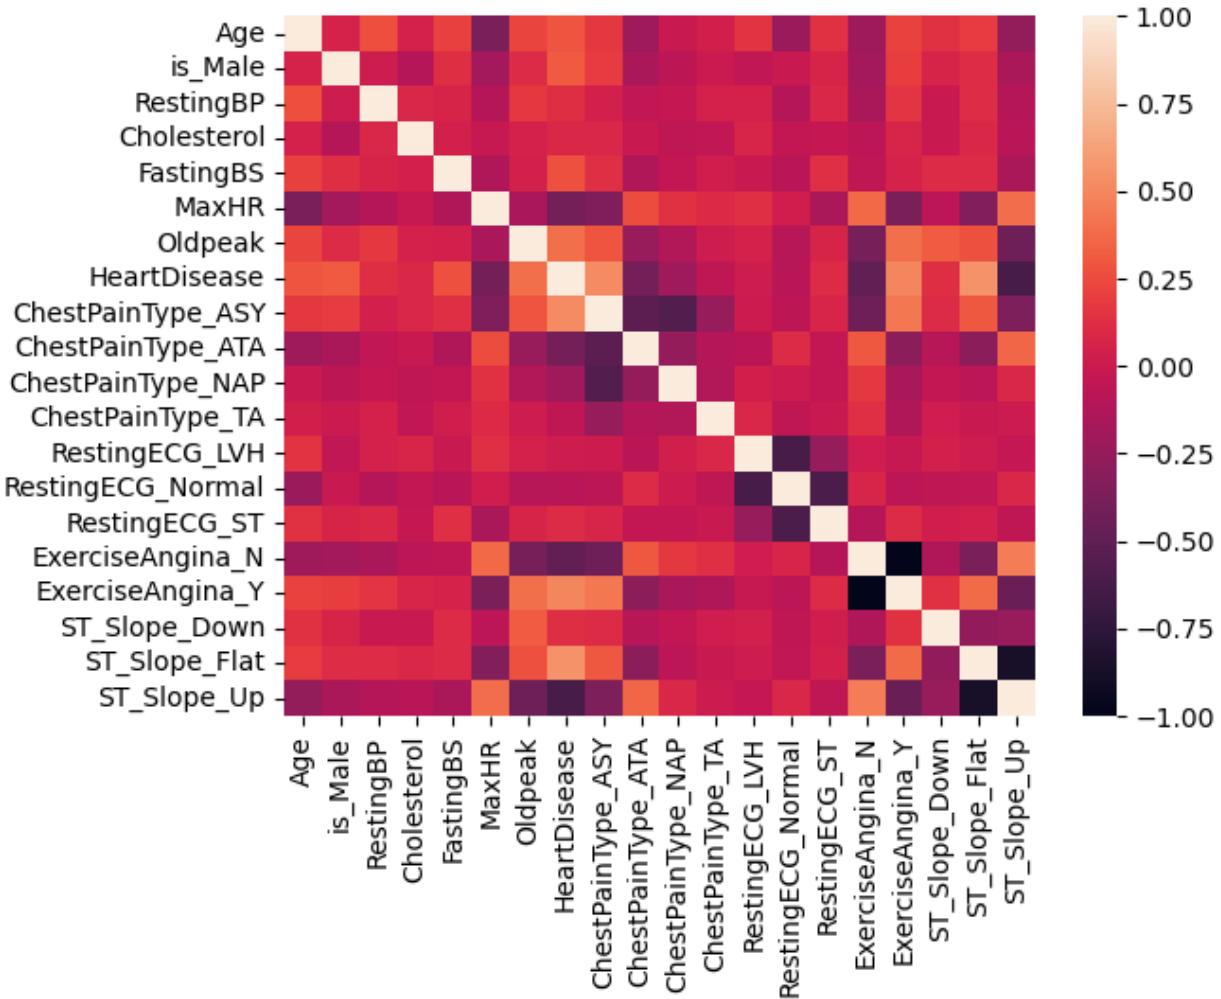




```
# from sklearn.preprocessing import StandardScaler
# cols=['Age', 'RestingBP', 'Cholesterol', 'MaxHR']
# scaler=StandardScaler()
# df_clean[cols]=scaler.fit_transform(df_clean[cols])
# df_clean

sns.heatmap(df_clean.corr())

<Axes: >
```



```

df_corr=df_clean.corr()['HeartDisease'].sort_values(ascending=False)

df_corr[df_corr>0]

HeartDisease      1.000000
ST_Slope_Flat     0.554134
ChestPainType_ASY 0.516716
ExerciseAngina_Y   0.494282
Oldpeak            0.392385
is_Male             0.305445
Age                0.282039
FastingBS          0.267291
ST_Slope_Down      0.122527
RestingBP           0.117909
RestingECG_ST       0.102527
Cholesterol         0.092586
RestingECG_LVH      0.010670
Name: HeartDisease, dtype: float64

```

```

df_clean.columns

Index(['Age', 'is_Male', 'RestingBP', 'Cholesterol', 'FastingBS',
'MaxHR',
       'Oldpeak', 'HeartDisease', 'ChestPainType_ASY',
'ChestPainType_ATA',
       'ChestPainType_NAP', 'ChestPainType_TA', 'RestingECG_LVH',
'RestingECG_Normal', 'RestingECG_ST', 'ExerciseAngina_N',
'ExerciseAngina_Y', 'ST_Slope_Down', 'ST_Slope_Flat',
'ST_Slope_Up'],
      dtype='object')

cat_features=[ 'is_Male', 'FastingBS','ChestPainType_ASY',
'ChestPainType_ATA', 'ChestPainType_NAP',
'ChestPainType_TA', 'RestingECG_LVH', 'RestingECG_Normal',
'RestingECG_ST','ExerciseAngina_N',
'ExerciseAngina_Y', 'ST_Slope_Down','ST_Slope_Flat',
'ST_Slope_Up']

from scipy.stats import chi2_contingency

alpha = 0.05

df_clean['HeartDisease_bin'] = df_clean['HeartDisease']
chi2_results = {}

for col in cat_features:
    contingency = pd.crosstab(df_clean[col],
df_clean['HeartDisease_bin'])
    chi2_stat, p_val, _, _ = chi2_contingency(contingency)
    decision = 'Reject Null (Keep Feature)' if p_val < alpha else
'Accept Null (Drop Feature)'
    chi2_results[col] = {
        'chi2_statistic': chi2_stat,
        'p_value': p_val,
        'Decision': decision
    }

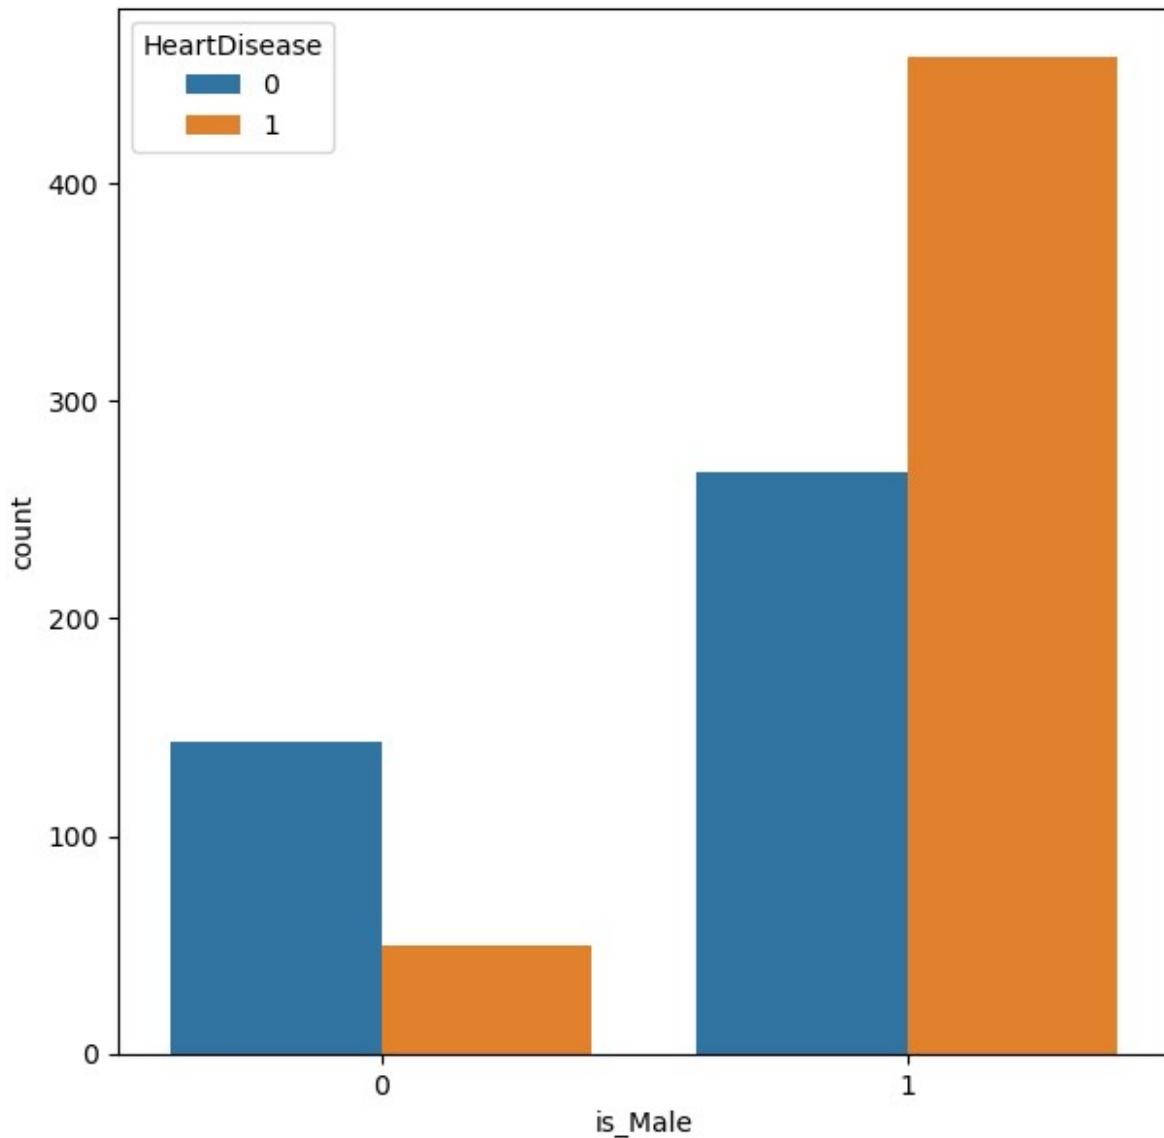
chi2_df = pd.DataFrame(chi2_results).T
chi2_df = chi2_df.sort_values(by='p_value')
chi2_df

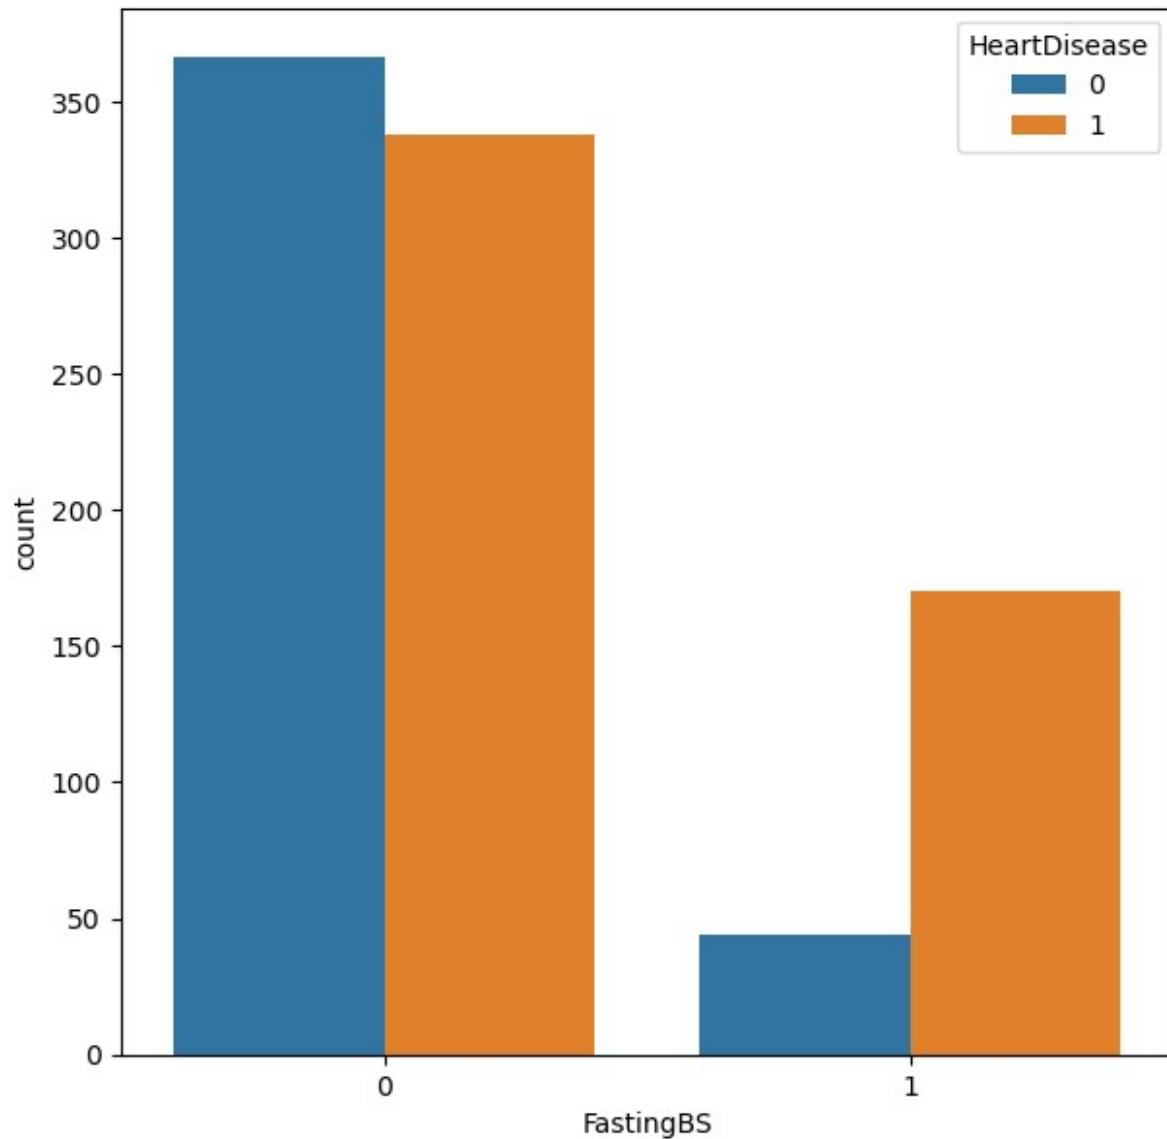
          chi2_statistic  p_value           Decision
ST_Slope_Up            352.823905     0.0  Reject Null (Keep Feature)
ST_Slope_Flat           279.659914     0.0  Reject Null (Keep Feature)
ChestPainType_ASY       243.021138     0.0  Reject Null (Keep Feature)
ExerciseAngina_N          222.259383     0.0  Reject Null (Keep Feature)
ExerciseAngina_Y          222.259383     0.0  Reject Null (Keep Feature)
ChestPainType_ATA        146.236323     0.0  Reject Null (Keep Feature)
is_Male                  84.145101     0.0  Reject Null (Keep Feature)
FastingBS                 64.320679     0.0  Reject Null (Keep Feature)

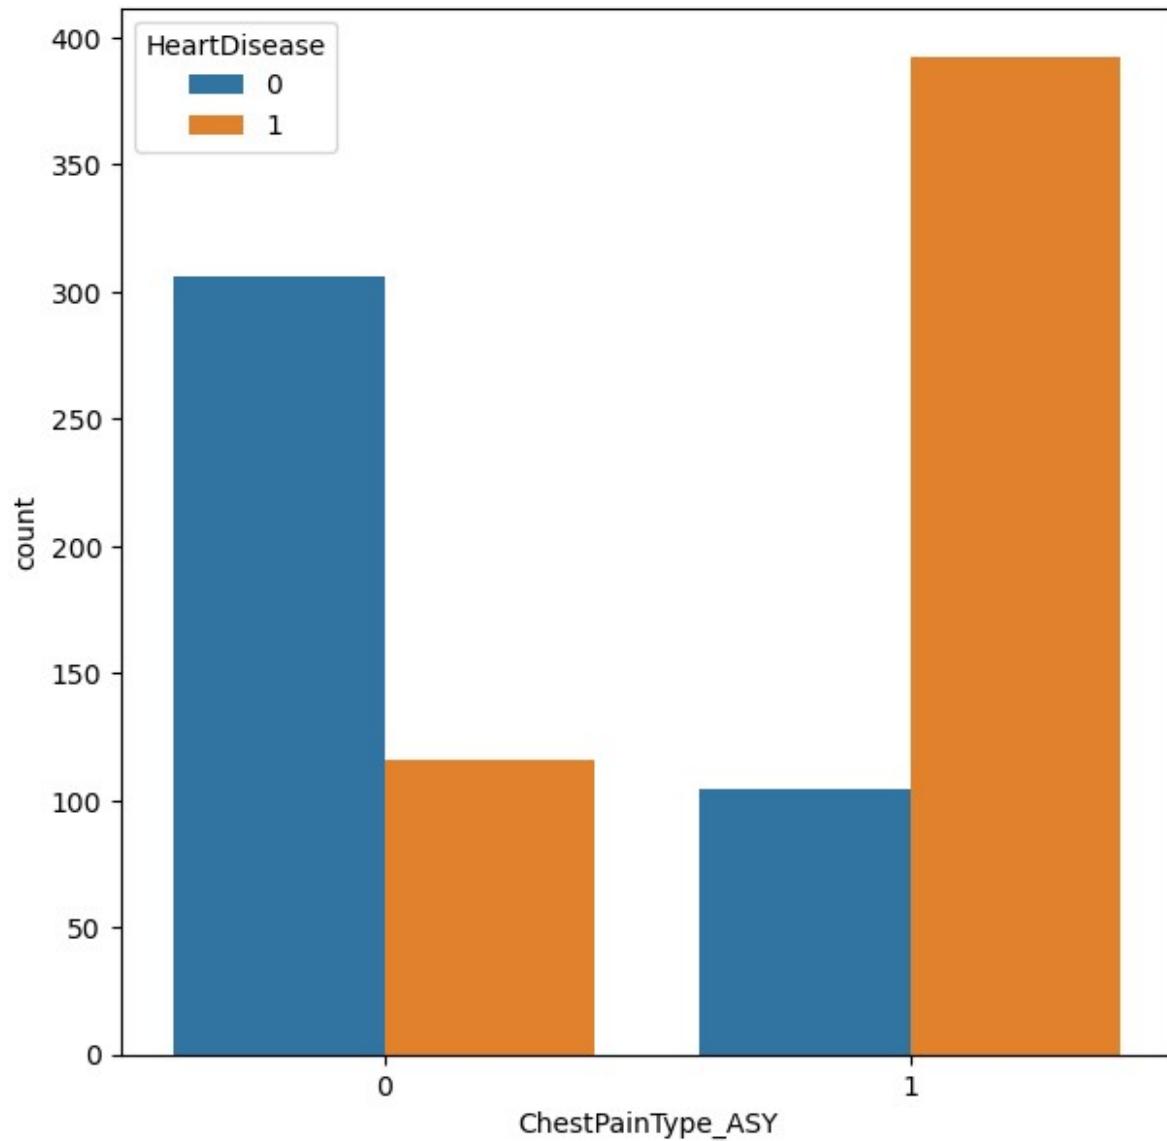
```

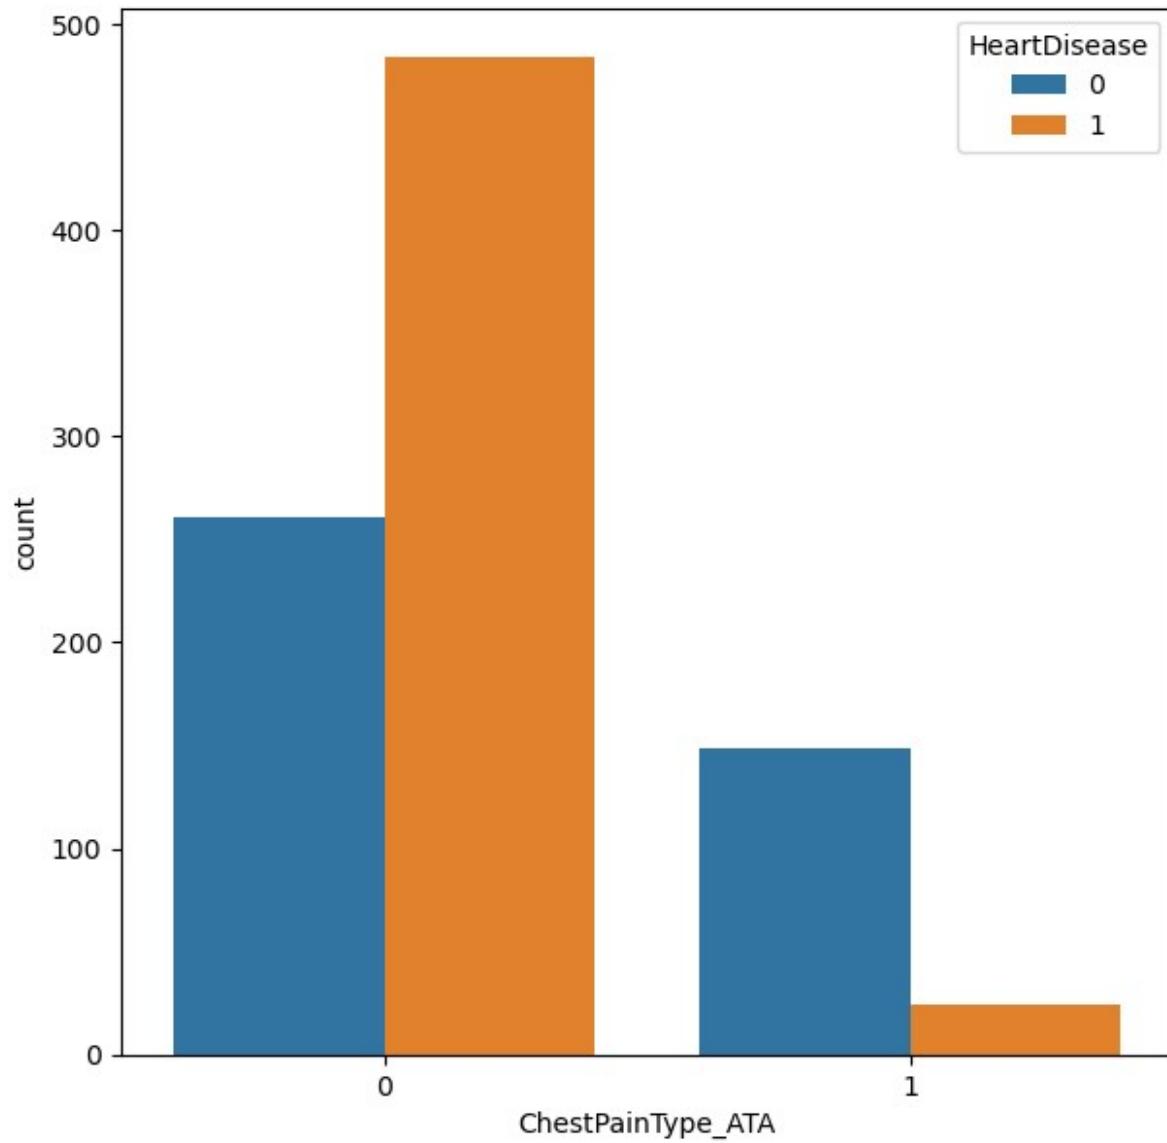
ChestPainType_NAP	40.608711	0.0	Reject	Null	(Keep Feature)
ST_Slope_Down	12.824125	0.000342	Reject	Null	(Keep Feature)
RestingECG_ST	9.135266	0.002507	Reject	Null	(Keep Feature)
RestingECG_Normal	7.327532	0.006791	Reject	Null	(Keep Feature)
ChestPainType_TA	2.273802	0.131577	Accept	Null	(Drop Feature)
RestingECG_LVH	0.058098	0.809528	Accept	Null	(Drop Feature)

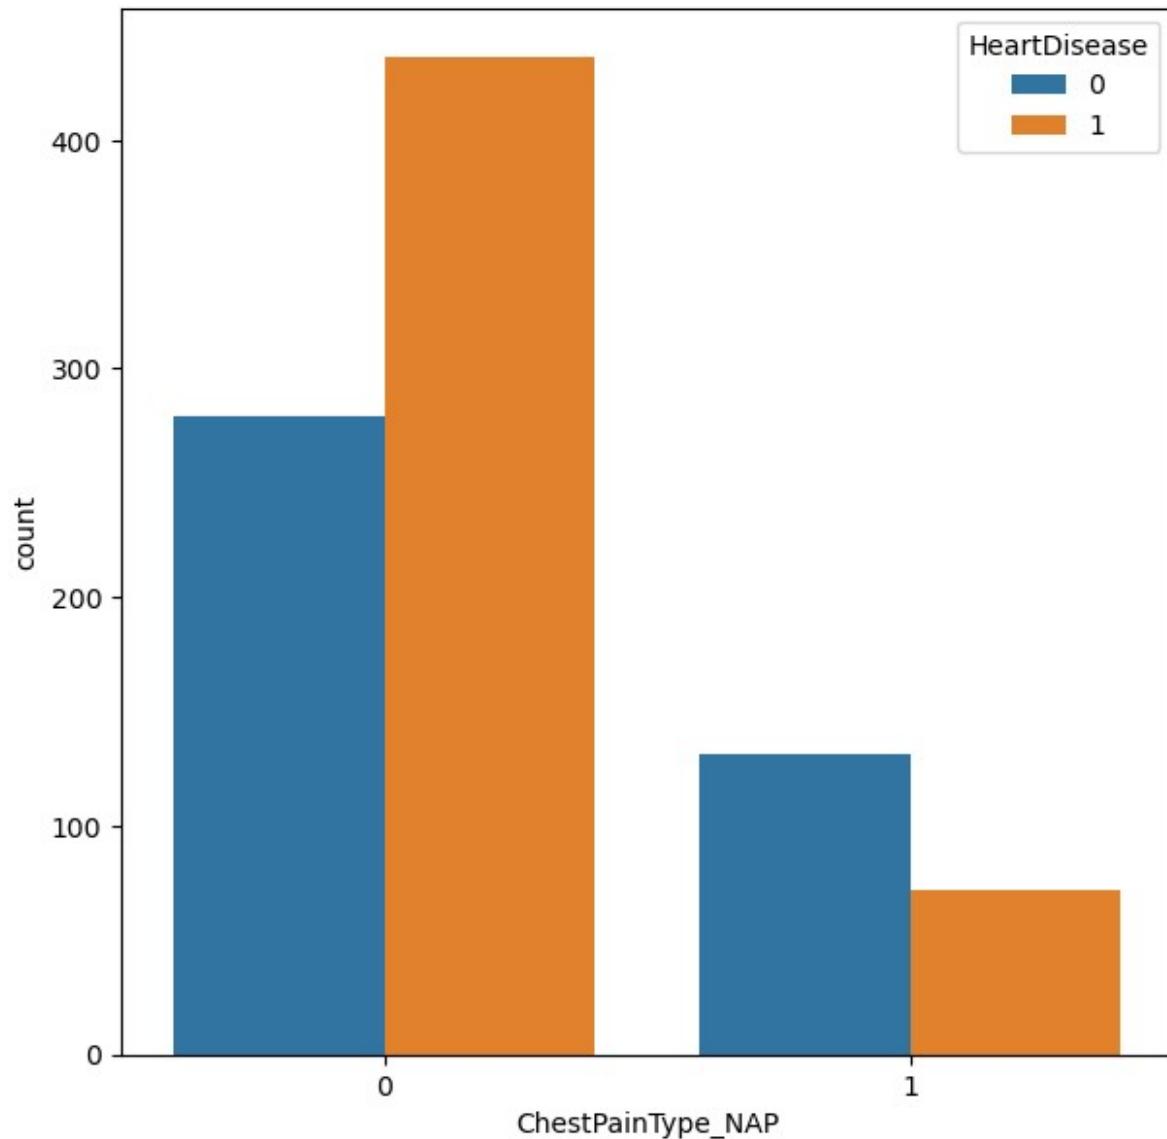
```
for i in cat_features:  
    plt.figure(figsize=(7,7))  
    sns.countplot(x=df_clean[i],hue=df_clean['HeartDisease'])
```

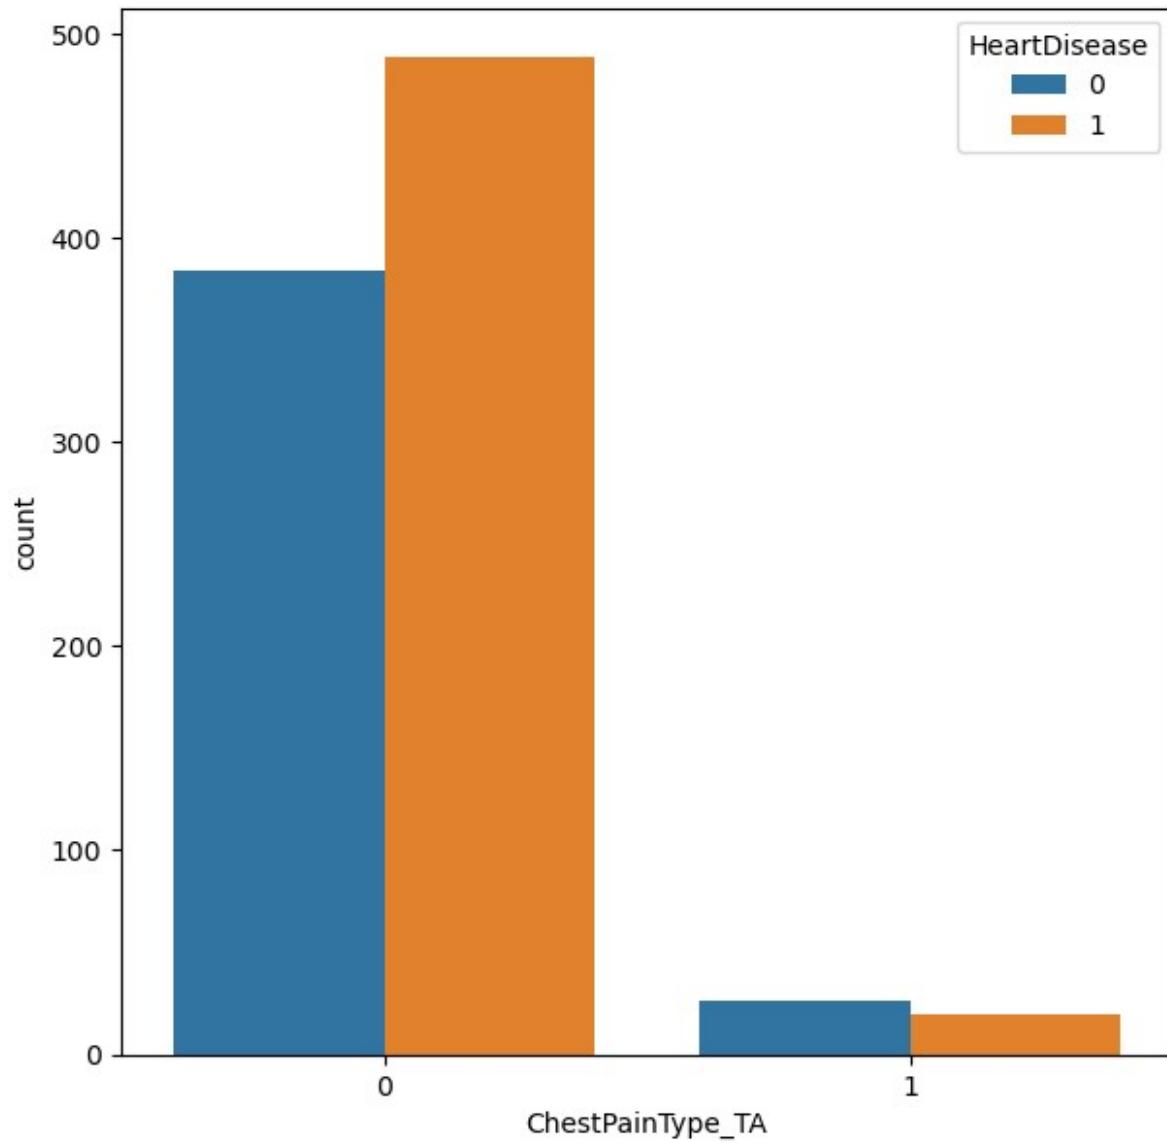


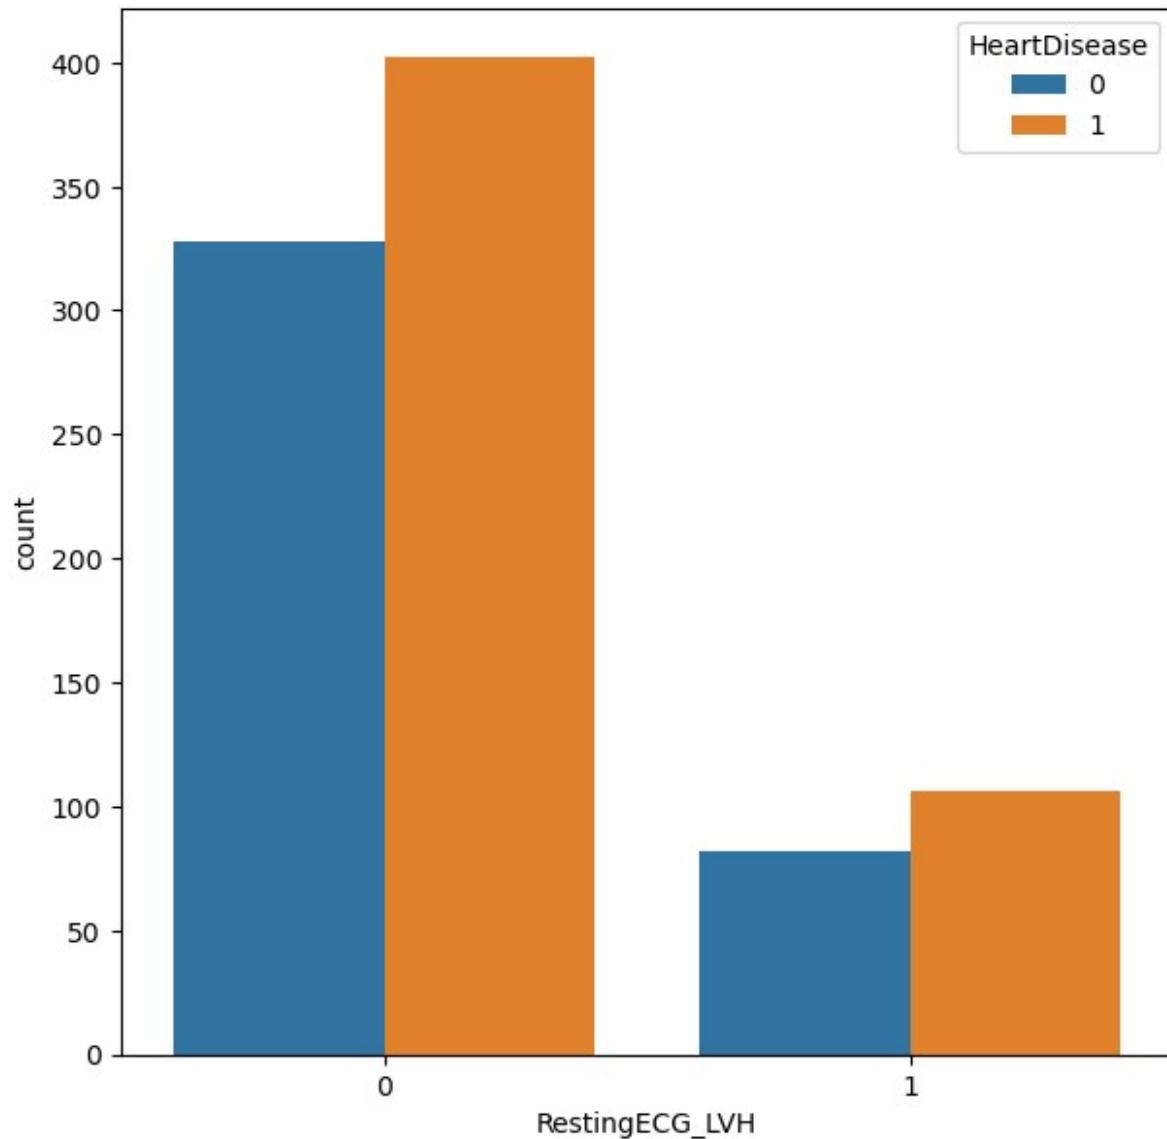


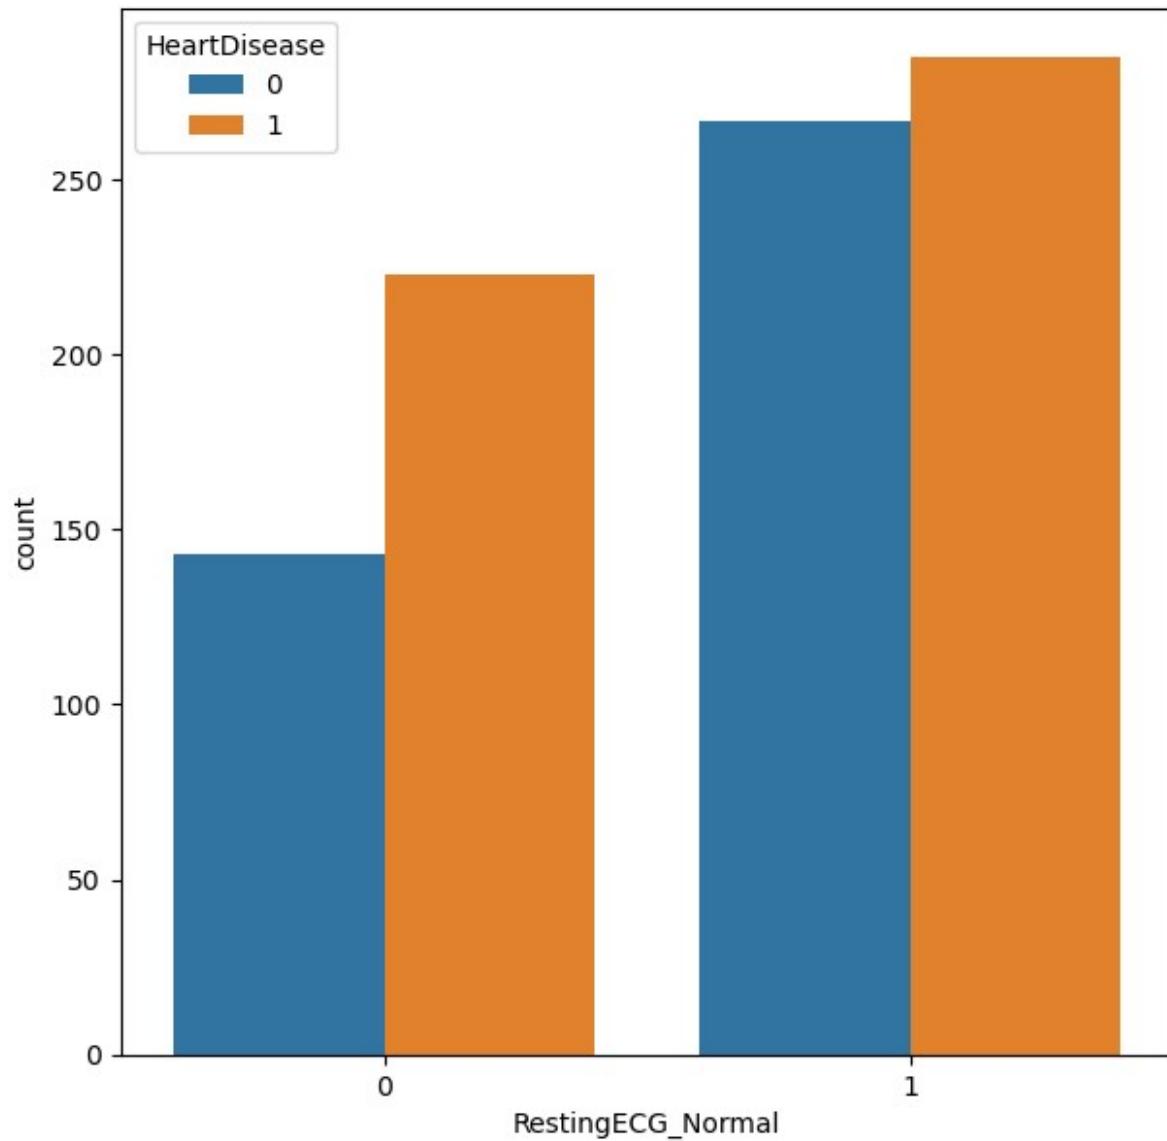


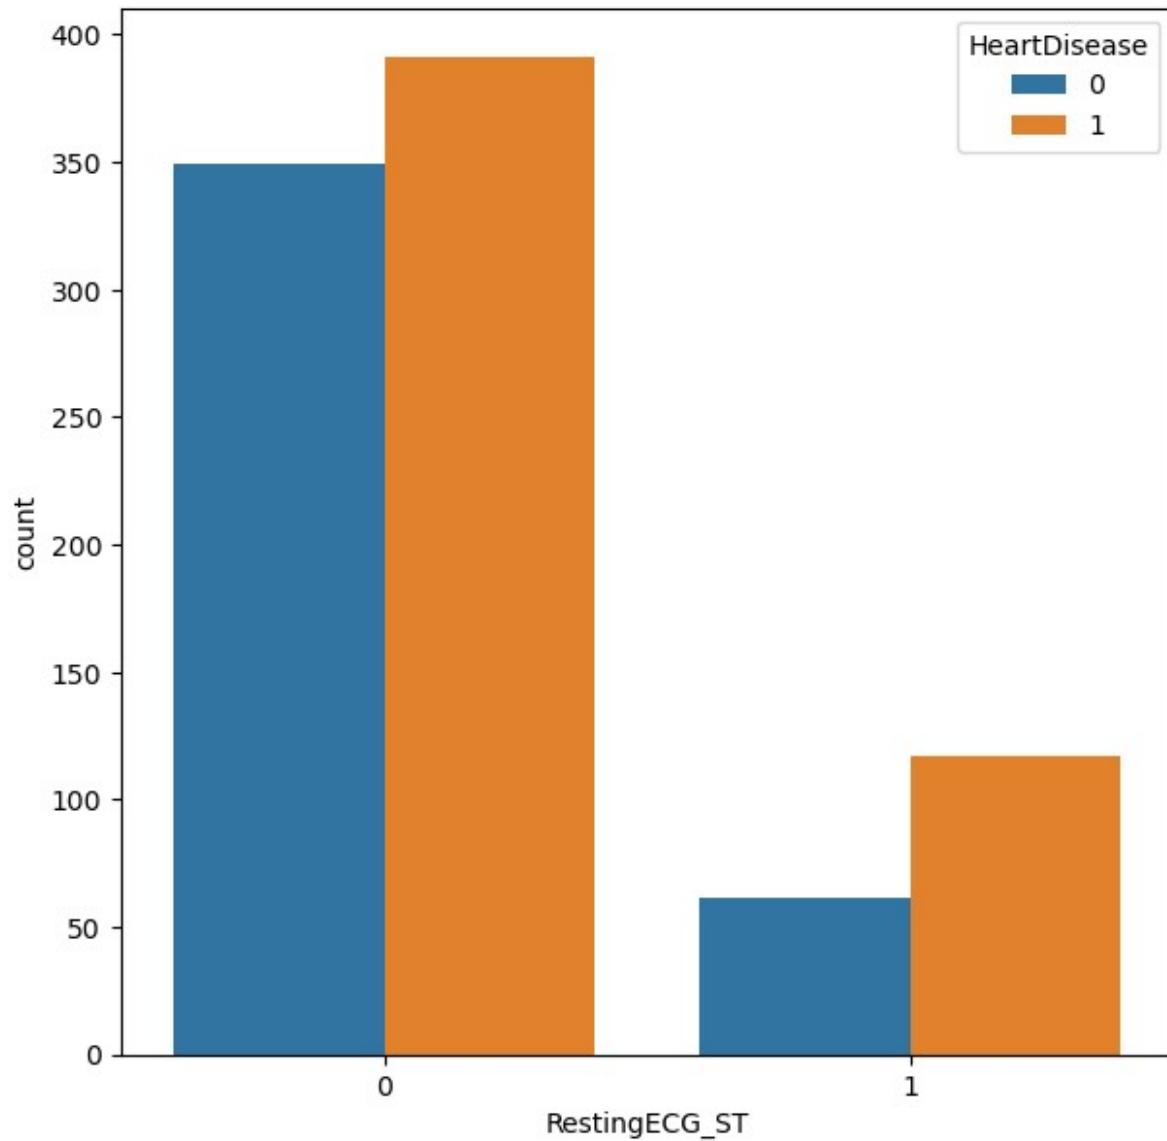


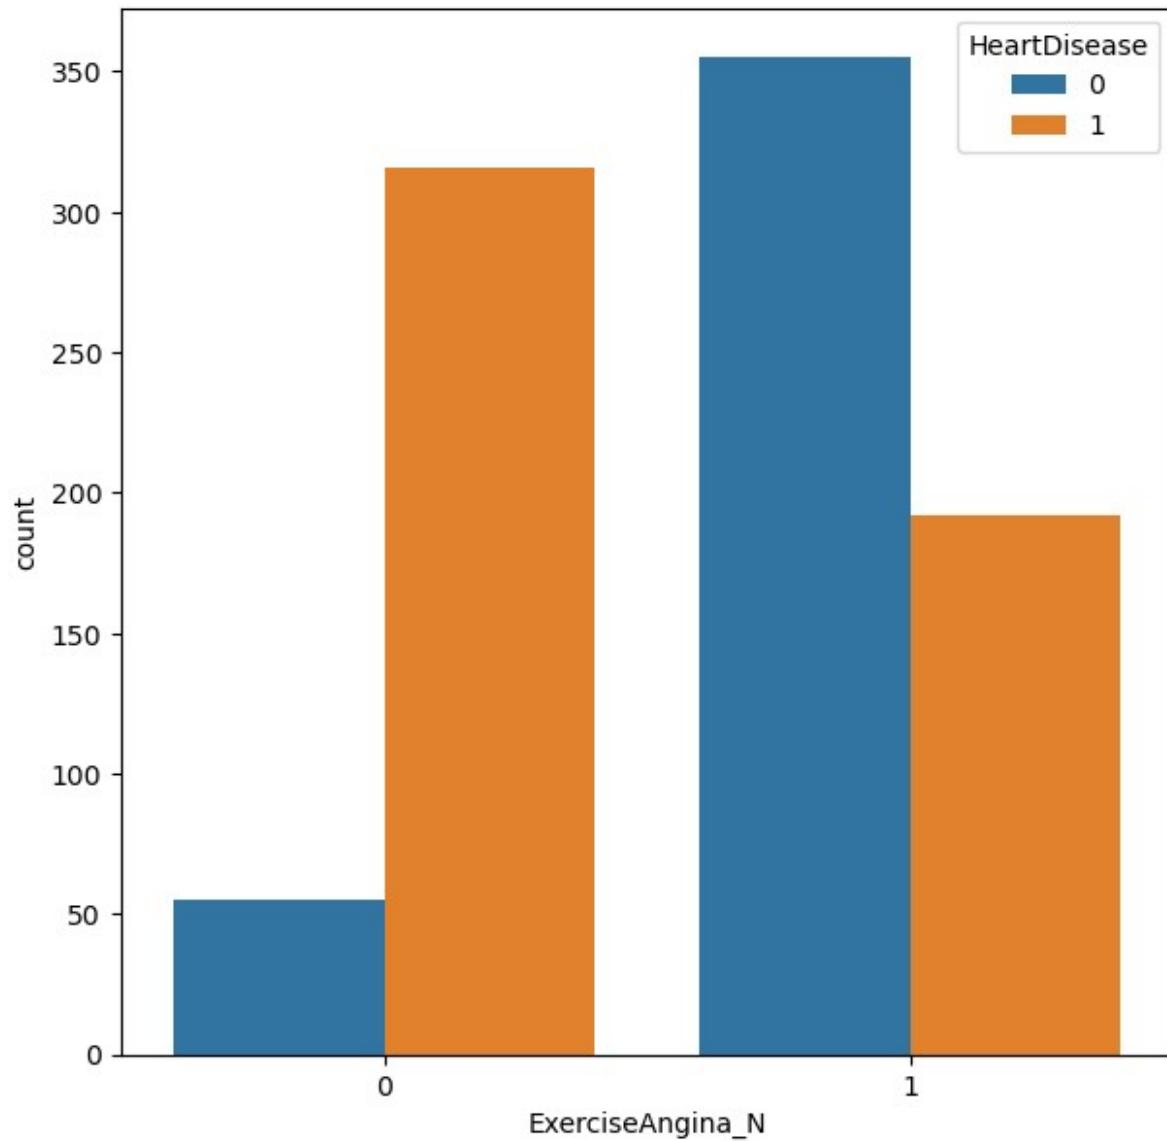


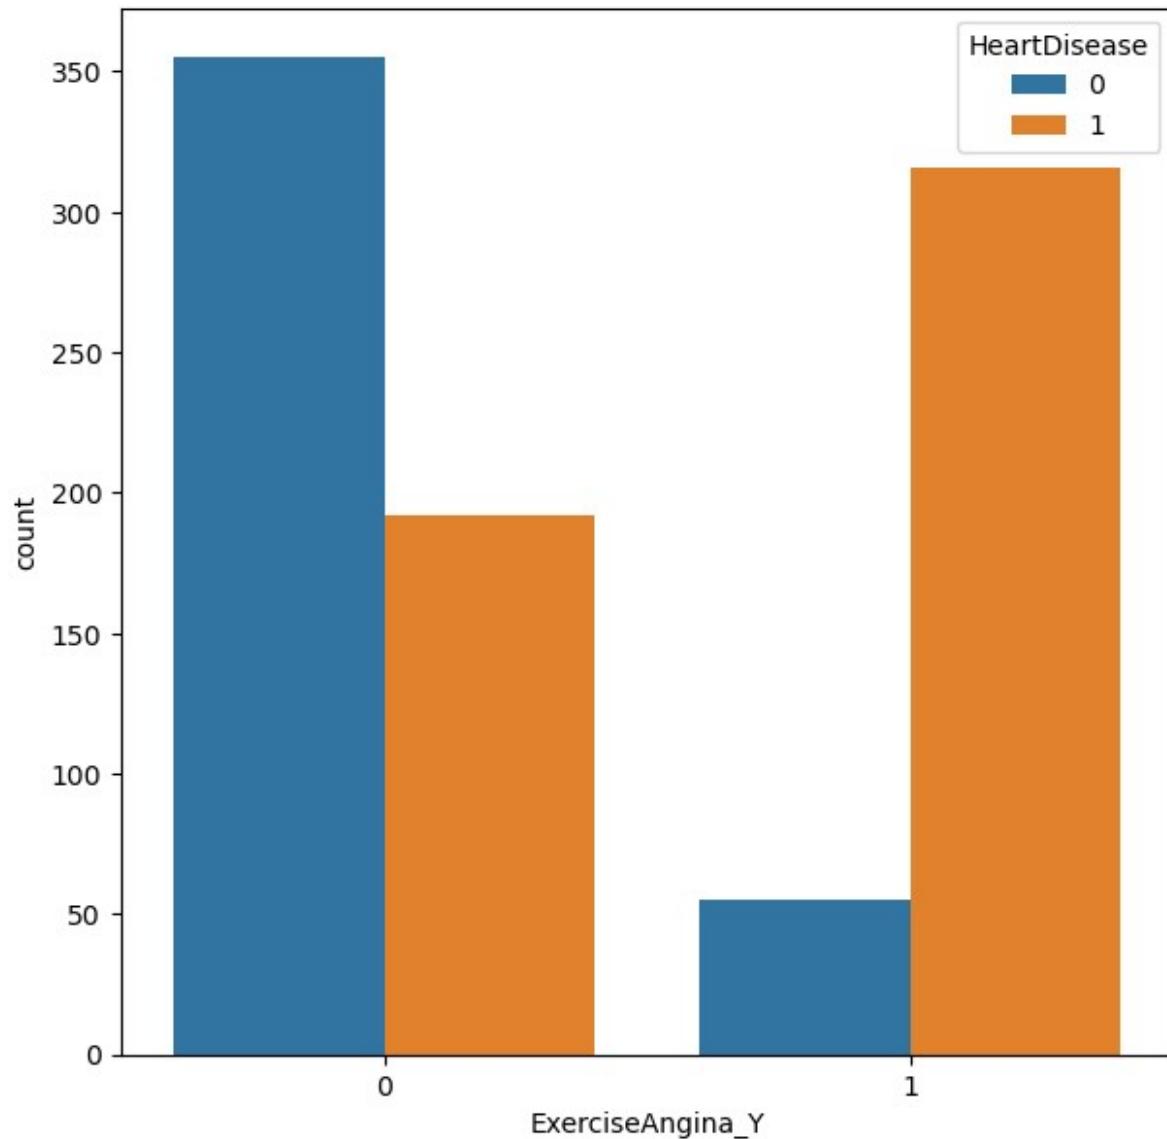


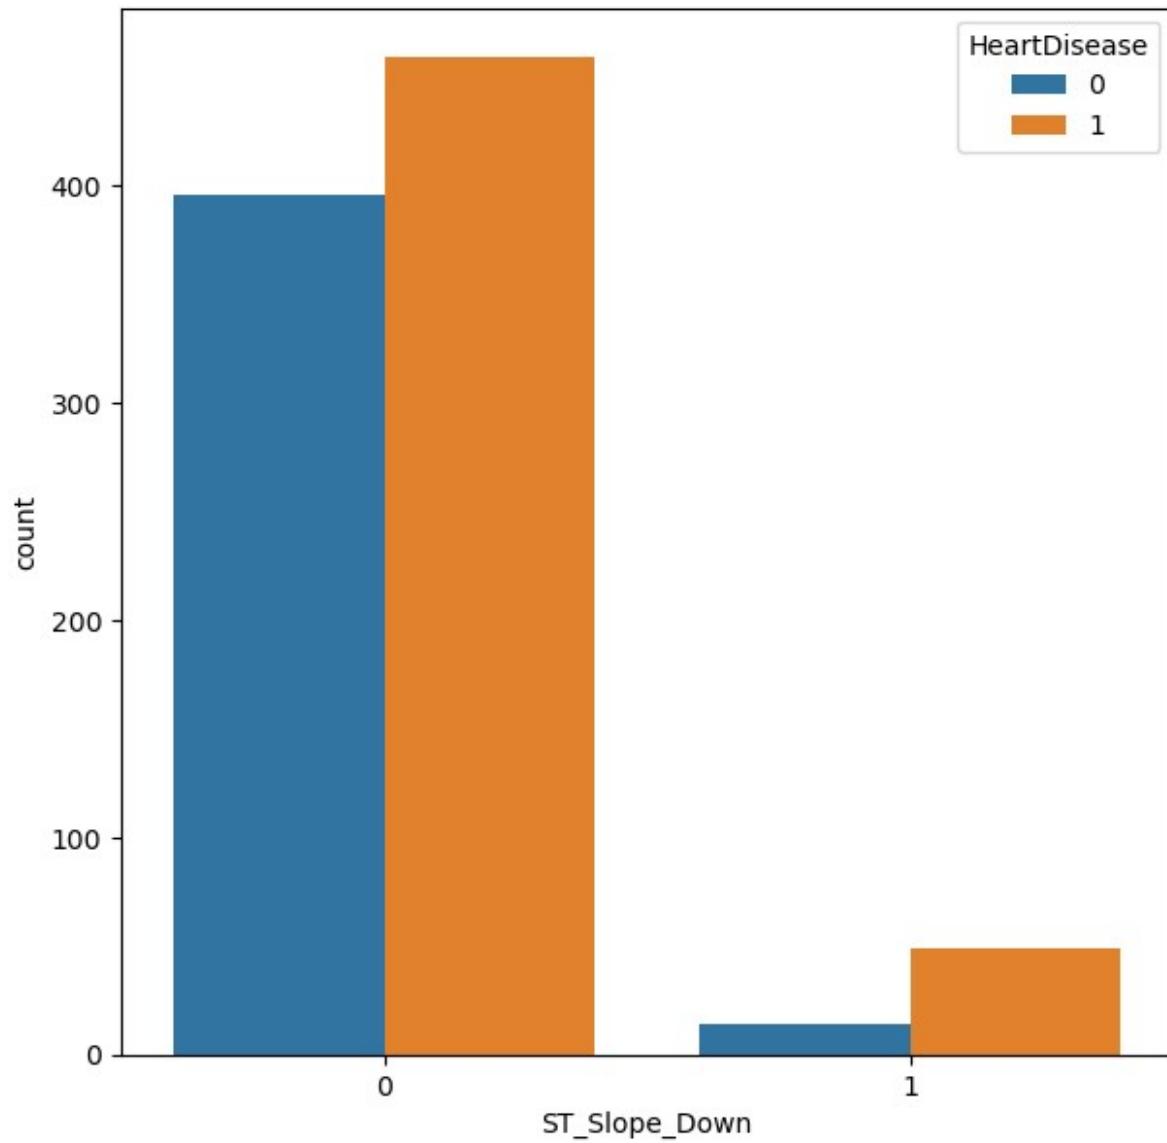


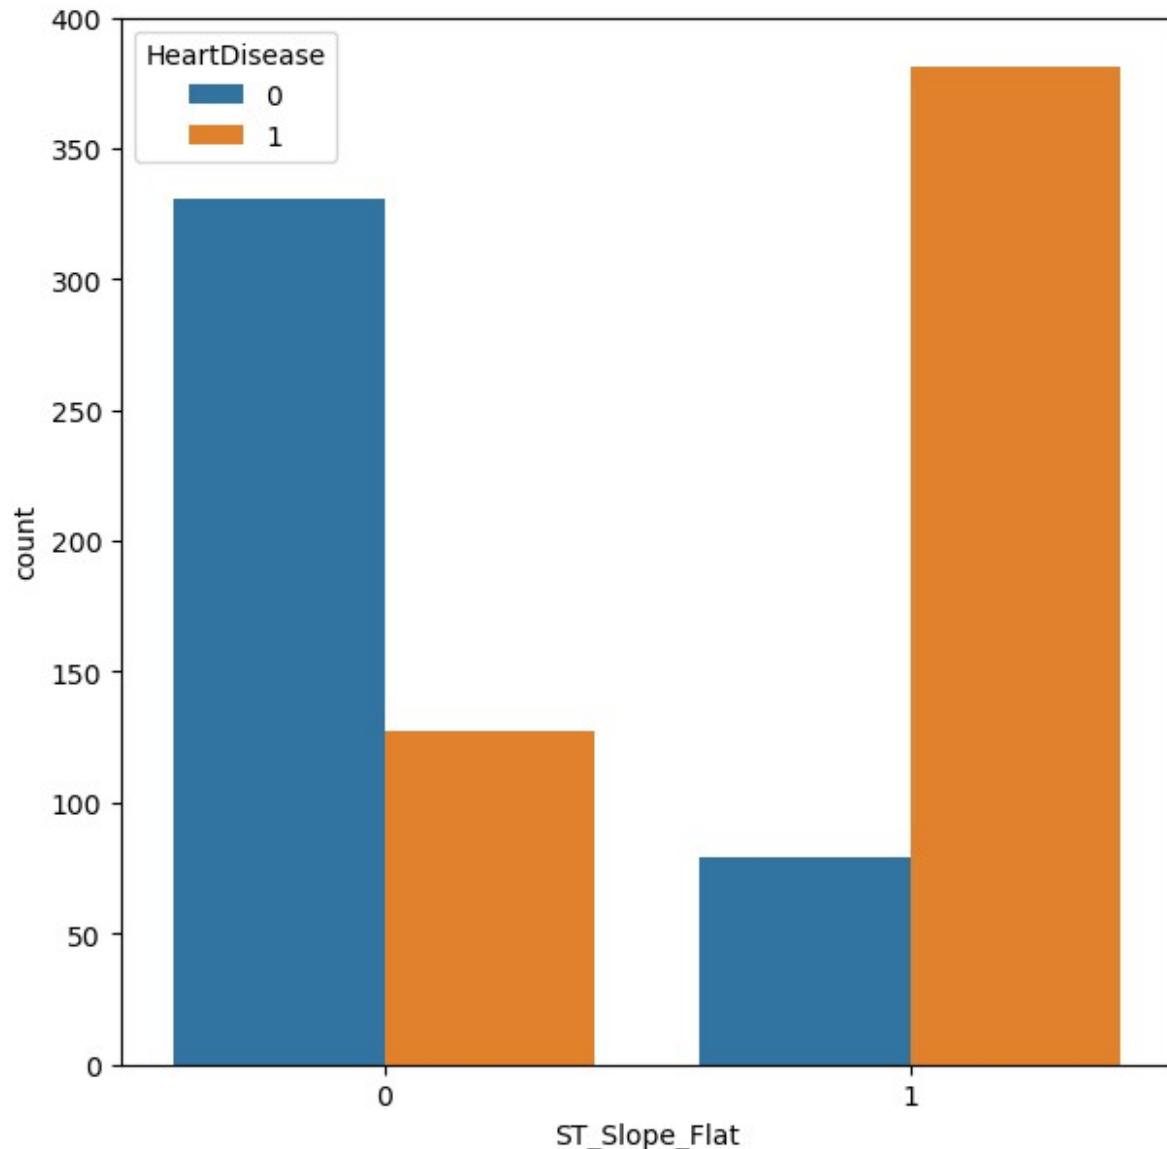


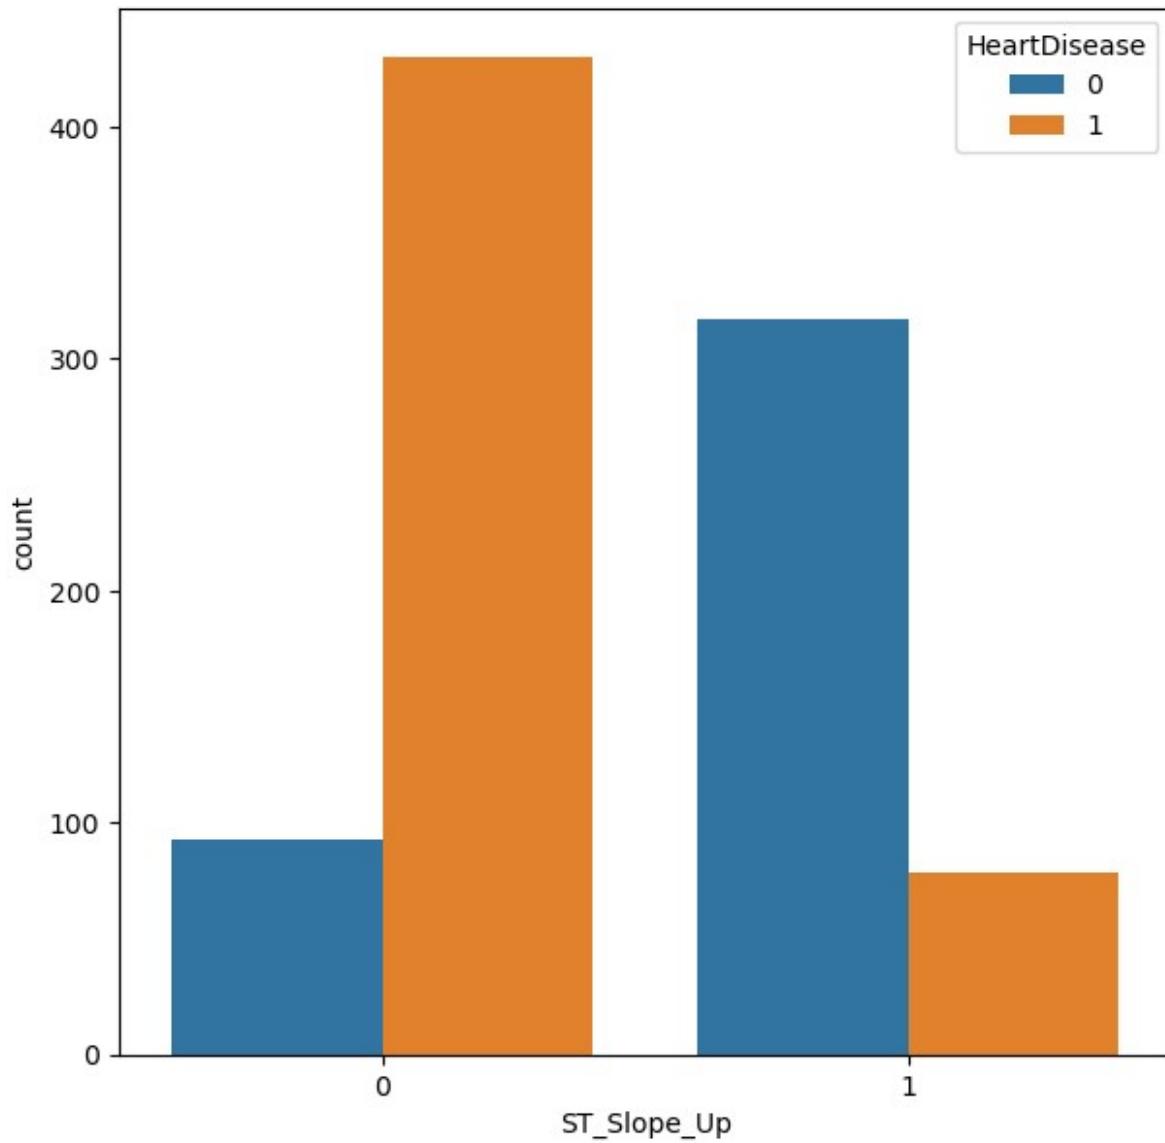












```
df_clean.columns  
Index(['Age', 'is_Male', 'RestingBP', 'Cholesterol', 'FastingBS',  
'MaxHR',  
       'Oldpeak', 'HeartDisease', 'ChestPainType_ASY',  
'ChestPainType_ATA',  
       'ChestPainType_NAP', 'ChestPainType_TA', 'RestingECG_LVH',  
       'RestingECG_Normal', 'RestingECG_ST', 'ExerciseAngina_N',  
       'ExerciseAngina_Y', 'ST_Slope_Down', 'ST_Slope_Flat',  
'ST_Slope_Up',  
       'HeartDisease_bin'],  
      dtype='object')  
df_clean.head(5)
```

	Age	is_Male	RestingBP	Cholesterol	FastingBS	MaxHR	Oldpeak	\
0	40	1	140	289	0	172	0	
1	49	0	160	180	0	156	1	
2	37	1	130	283	0	98	0	
3	48	0	138	214	0	108	1	
4	54	1	150	195	0	122	0	
HeartDisease ChestPainType_ASY ChestPainType_ATA ...								
ChestPainType_TA \								
0		0		0		1	...	
0								
1		1		0		0	...	
0								
2		0		0		1	...	
0								
3		1		1		0	...	
0								
4		0		0		0	...	
0								
RestingECG_LVH RestingECG_Normal RestingECG_ST ExerciseAngina_N								
\								
0		0		1		0		1
1		0		1		0		1
2		0		0		1		1
3		0		1		0		0
4		0		1		0		1
ExerciseAngina_Y ST_Slope_Down ST_Slope_Flat ST_Slope_Up \								
0		0		0		0		1
1		0		0		1		0
2		0		0		0		1
3		1		0		1		0
4		0		0		0		1
HeartDisease_bin								
0		0						
1		1						
2		0						
3		1						
4		0						

model

```
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import accuracy_score, f1_score,
classification_report
from sklearn.linear_model import LogisticRegression
from sklearn.naive_bayes import GaussianNB
from sklearn.tree import DecisionTreeClassifier
from sklearn.svm import SVC
from sklearn.neighbors import KNeighborsClassifier

df_clean.columns
Index(['Age', 'is_Male', 'RestingBP', 'Cholesterol', 'FastingBS',
'MaxHR',
       'Oldpeak', 'HeartDisease', 'ChestPainType_ASY',
'ChestPainType_ATA',
       'ChestPainType_NAP', 'ChestPainType_TA', 'RestingECG_LVH',
'RestingECG_Normal', 'RestingECG_ST', 'ExerciseAngina_N',
'ExerciseAngina_Y', 'ST_Slope_Down', 'ST_Slope_Flat',
'ST_Slope_Up',
       'HeartDisease_bin'],
      dtype='object')

X = df_clean.drop(columns=[ 'HeartDisease', 'ExerciseAngina_N',
                           'HeartDisease_bin'], axis=1)
y = df_clean['HeartDisease']

X_train, X_test, y_train, y_test = train_test_split(
    X, y, stratify=y, test_size=0.2, random_state=42)

scaler=StandardScaler()
x_train_scaled=scaler.fit_transform(X_train)
x_test_scaled=scaler.fit_transform(X_test)

models = {
    "Logistic Regression": LogisticRegression(),
    "KNN": KNeighborsClassifier(n_neighbors=7),
    "Naive Bayes": GaussianNB(),
    "Decision Tree": DecisionTreeClassifier(),
    "SVM (RBF Kernel)": SVC(probability=True,kernel='poly')
}

results = []

for name, model in models.items():
    model.fit(x_train_scaled, y_train)
    y_pred = model.predict(x_test_scaled)
    acc = accuracy_score(y_test, y_pred)
    f1 = f1_score(y_test, y_pred)
```

```
results.append({
    'Model': name,
    'Accuracy': round(acc, 4),
    'F1 Score': round(f1, 4)
})

results
[{'Model': 'Logistic Regression', 'Accuracy': 0.8967, 'F1 Score': 0.91},
 {'Model': 'KNN', 'Accuracy': 0.875, 'F1 Score': 0.89},
 {'Model': 'Naive Bayes', 'Accuracy': 0.8587, 'F1 Score': 0.8725},
 {'Model': 'Decision Tree', 'Accuracy': 0.7554, 'F1 Score': 0.7739},
 {'Model': 'SVM (RBF Kernel)', 'Accuracy': 0.875, 'F1 Score': 0.891}]

import joblib
joblib.dump(models['Logistic Regression'],'LOG_heart.pkl')
joblib.dump(scaler,'scaler.pkl')
joblib.dump(X.columns.tolist(),'columns.pkl')

['columns.pkl']

X.columns
Index(['Age', 'is_Male', 'RestingBP', 'Cholesterol', 'FastingBS',
       'MaxHR',
       'Oldpeak', 'ChestPainType_ASY', 'ChestPainType_ATA',
       'ChestPainType_NAP', 'ChestPainType_TA', 'RestingECG_LVH',
       'RestingECG_Normal', 'RestingECG_ST', 'ExerciseAngina_Y',
       'ST_Slope_Down', 'ST_Slope_Flat', 'ST_Slope_Up'],
      dtype='object')
```